



# Convolutional Neural Networks

Amir H. Payberah  
[payberah@kth.se](mailto:payberah@kth.se)  
05/12/2018



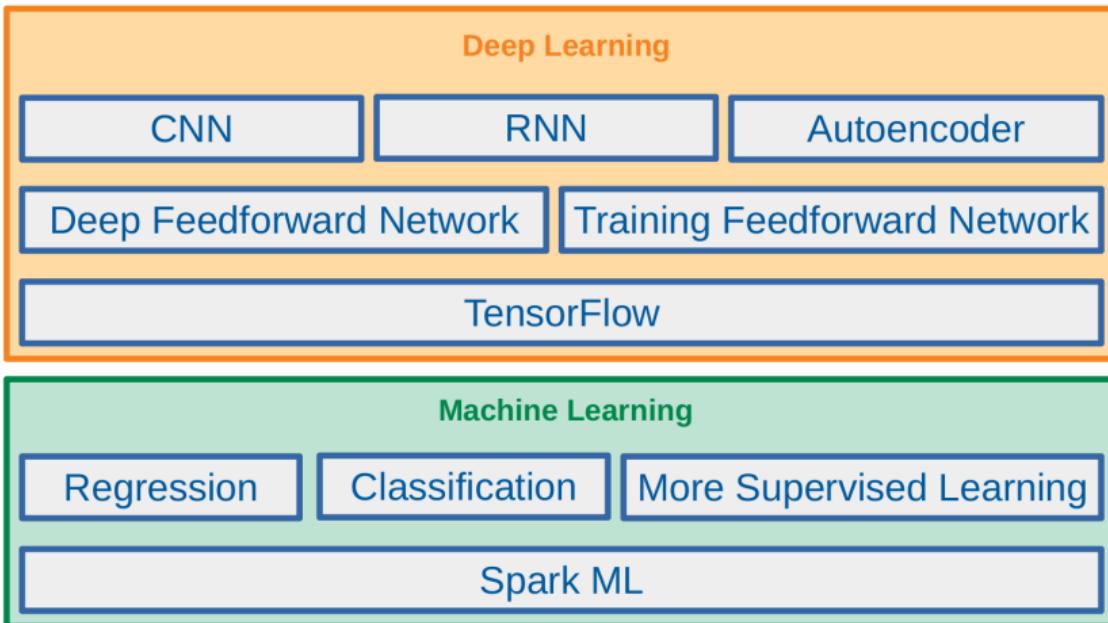


# The Course Web Page

<https://id2223kth.github.io>

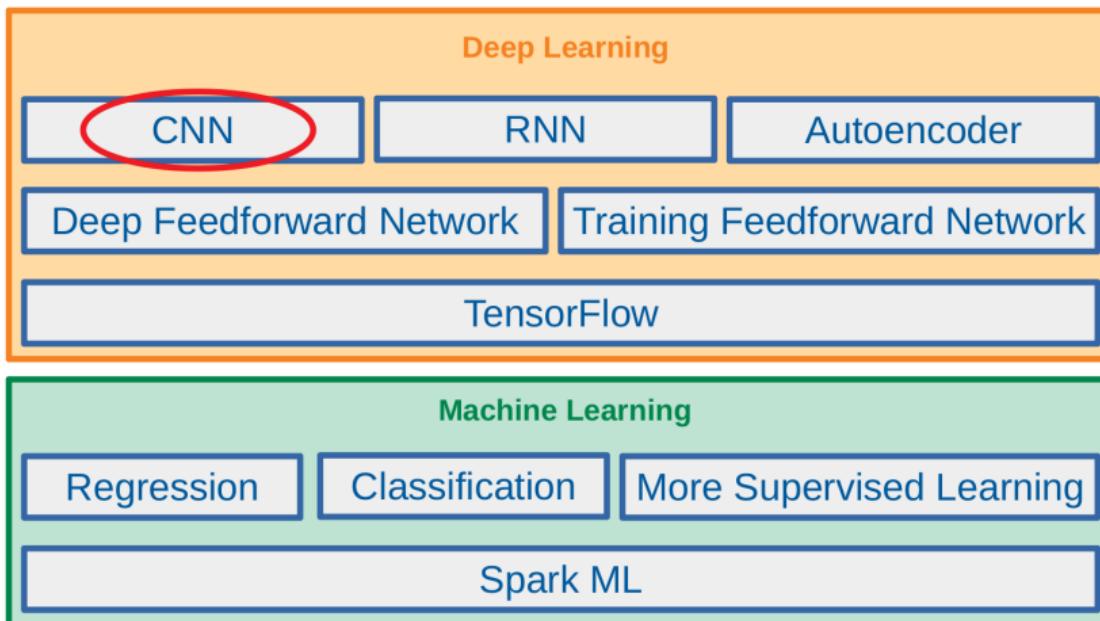


# Where Are We?





# Where Are We?





# Let's Start With An Example



## MNIST Dataset

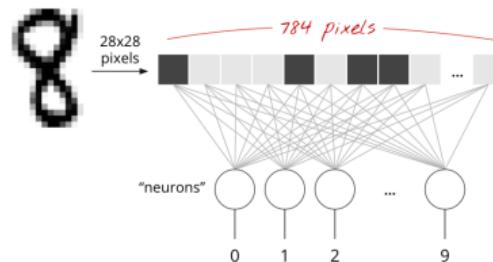
- ▶ Handwritten digits in the [MNIST](#) dataset are 28x28 pixel greyscale images.

A grid of handwritten digits from the MNIST dataset, arranged in ten rows. Each row contains ten digits of a specific type, demonstrating the variety of handwriting styles in the dataset. The digits are 28x28 pixel greyscale images.

0	0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9	9

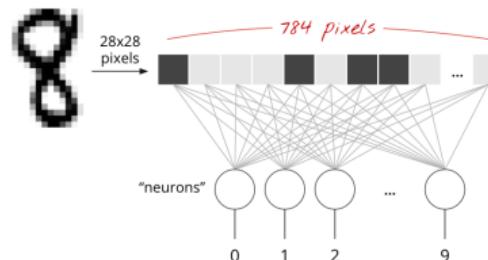
# One-Layer Network For Classifying MNIST (1/4)

- ▶ Let's make a **one-layer** neural network for **classifying digits**.



# One-Layer Network For Classifying MNIST (1/4)

- ▶ Let's make a **one-layer** neural network for **classifying digits**.
- ▶ Each **neuron** in a neural network:
  - Does a **weighted sum** of all of its inputs
  - Adds a **bias**
  - Feeds the result through some **non-linear activation** function, e.g., **softmax**.



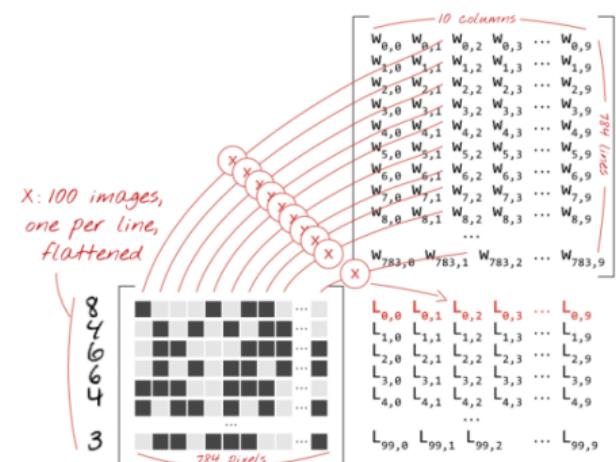
# One-Layer Network For Classifying MNIST (2/4)



[<https://github.com/GoogleCloudPlatform/tensorflow-without-a-phd>]

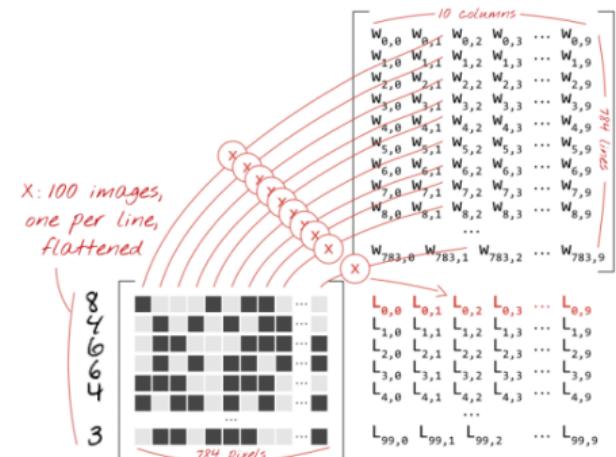
# One-Layer Network For Classifying MNIST (3/4)

- ▶ Assume we have a **batch of 100 images** as the **input**.



# One-Layer Network For Classifying MNIST (3/4)

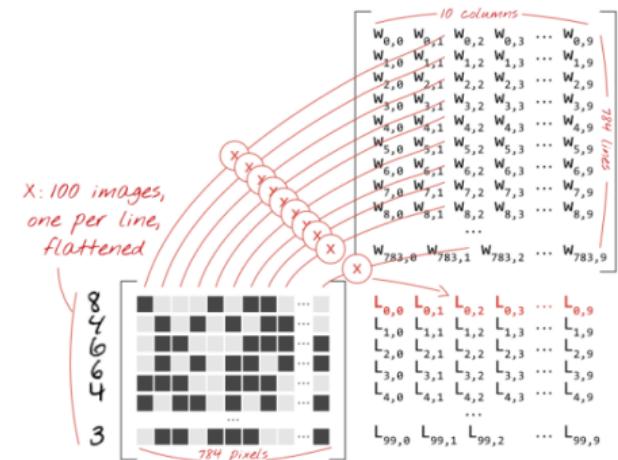
- ▶ Assume we have a **batch of 100 images** as the **input**.
- ▶ Using the **first column** of the **weights matrix  $\mathbf{W}$** , we compute the **weighted sum** of all the **pixels** of the **first image**.



# One-Layer Network For Classifying MNIST (3/4)

- ▶ Assume we have a **batch of 100 images** as the **input**.
- ▶ Using the **first column** of the **weights matrix  $\mathbf{W}$** , we compute the **weighted sum** of all the **pixels** of the **first image**.
  - The **first neuron**:

$$L_{0,0} = w_{0,0}x_0^{(1)} + w_{1,0}x_1^{(1)} + \dots + w_{783,0}x_{783}^{(1)}$$



# One-Layer Network For Classifying MNIST (3/4)

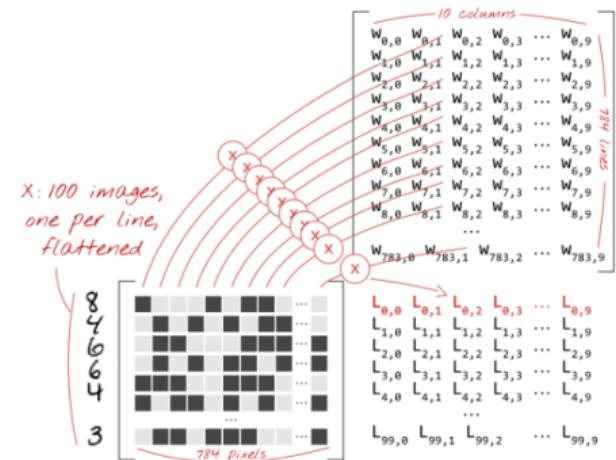
- ▶ Assume we have a **batch of 100 images** as the **input**.
- ▶ Using the **first column** of the **weights matrix  $\mathbf{W}$** , we compute the **weighted sum** of all the **pixels** of the **first image**.
  - The **first neuron**:  

$$L_{0,0} = w_{0,0}x_0^{(1)} + w_{1,0}x_1^{(1)} + \dots + w_{783,0}x_{783}^{(1)}$$
  - The **2nd neuron until the 10th**:  

$$L_{0,1} = w_{0,1}x_0^{(1)} + w_{1,1}x_1^{(1)} + \dots + w_{783,1}x_{783}^{(1)}$$

$$\dots$$

$$L_{0,9} = w_{0,9}x_0^{(1)} + w_{1,9}x_1^{(1)} + \dots + w_{783,9}x_{783}^{(1)}$$



# One-Layer Network For Classifying MNIST (3/4)

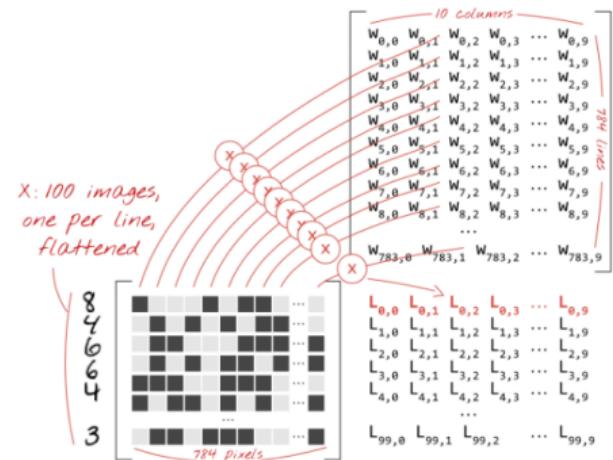
- ▶ Assume we have a **batch of 100 images** as the **input**.
- ▶ Using the **first column** of the **weights matrix  $\mathbf{W}$** , we compute the **weighted sum** of all the **pixels** of the **first image**.
  - The **first neuron**:  

$$L_{0,0} = w_{0,0}x_0^{(1)} + w_{1,0}x_1^{(1)} + \dots + w_{783,0}x_{783}^{(1)}$$
  - The **2nd neuron until the 10th**:  

$$L_{0,1} = w_{0,1}x_0^{(1)} + w_{1,1}x_1^{(1)} + \dots + w_{783,1}x_{783}^{(1)}$$

$$\dots$$

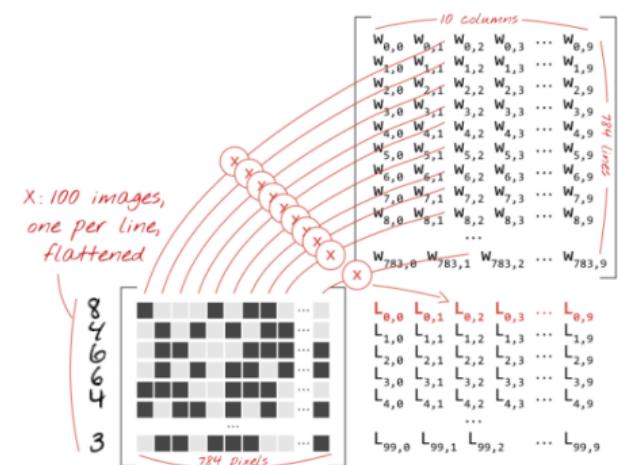
$$L_{0,9} = w_{0,9}x_0^{(1)} + w_{1,9}x_1^{(1)} + \dots + w_{783,9}x_{783}^{(1)}$$
  - Repeat the operation for the **other 99 images**, i.e.,  $\mathbf{x}^{(2)} \dots \mathbf{x}^{(100)}$



# One-Layer Network For Classifying MNIST (4/4)

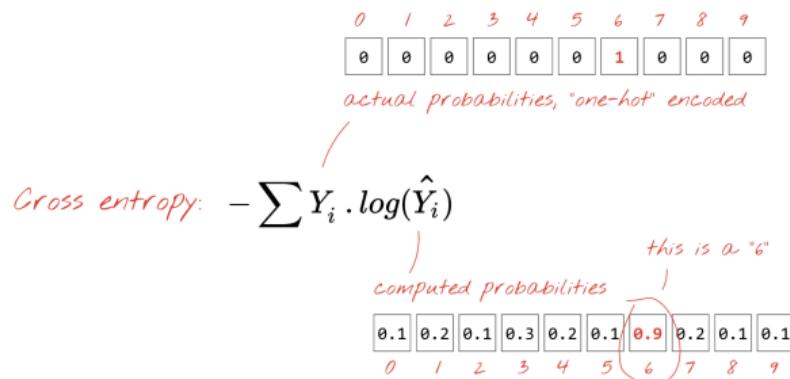
- ▶ Each neuron must now add its **bias**.
- ▶ Apply the **softmax activation function** for each instance  $\mathbf{x}^{(i)}$ .

- ▶ For each input instance  $\mathbf{x}^{(i)}$ :  $\mathbf{L}_i = \begin{bmatrix} L_{i,0} \\ L_{i,1} \\ \vdots \\ L_{i,9} \end{bmatrix}$
- ▶  $\hat{\mathbf{y}}_i = \text{softmax}(\mathbf{L}_i + \mathbf{b})$



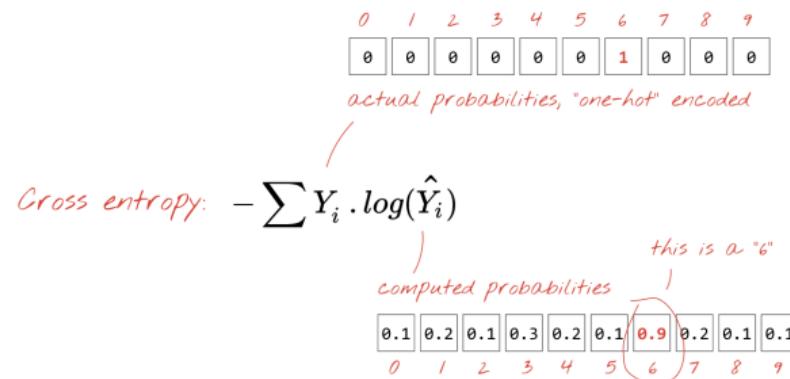
# How Good the Predictions Are?

- ▶ Define the cost function  $J(\mathbf{W})$  as the **cross-entropy** of what the network tells us  $\hat{\mathbf{y}}_i$  and what we know to be the truth  $\mathbf{y}_i$ , for each instance  $\mathbf{x}^{(i)}$ .



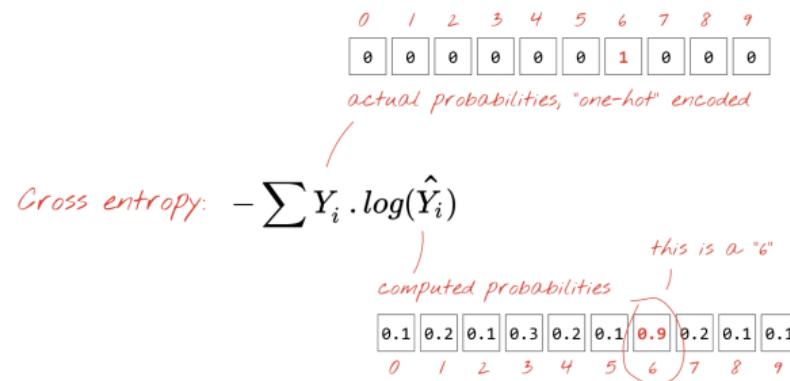
# How Good the Predictions Are?

- ▶ Define the cost function  $J(\mathbf{W})$  as the **cross-entropy** of what the network tells us  $\hat{\mathbf{y}}_i$  and what we know to be the truth  $\mathbf{y}_i$ , for each instance  $\mathbf{x}^{(i)}$ .
- ▶ Compute the **partial derivatives** of the cross-entropy relatively to all the **weights** and **all the biases**,  $\nabla_{\mathbf{W}} J(\mathbf{W})$ .



# How Good the Predictions Are?

- ▶ Define the cost function  $J(\mathbf{W})$  as the **cross-entropy** of what the network tells us  $\hat{\mathbf{y}}_i$  and what we know to be the truth  $\mathbf{y}_i$ , for each instance  $\mathbf{x}^{(i)}$ .
- ▶ Compute the **partial derivatives** of the cross-entropy relatively to all the **weights** and all the **biases**,  $\nabla_{\mathbf{W}} J(\mathbf{W})$ .
- ▶ Update weights and biases by a **fraction of the gradient**  $\mathbf{W}^{(\text{next})} = \mathbf{W} - \eta \nabla_{\mathbf{W}} J(\mathbf{W})$

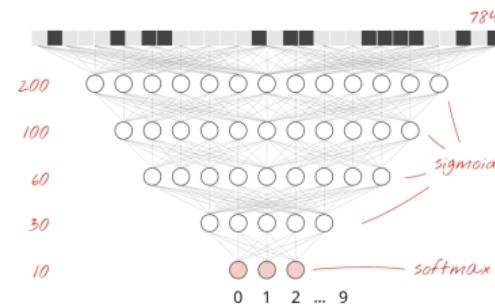


# Adding More Layers

- ▶ Add more layers to **improve the accuracy**.
- ▶ On **intermediate layers** we will use the the **sigmoid** activation function.
- ▶ We keep **softmax** as the activation function on the **last layer**.



[<https://github.com/GoogleCloudPlatform/tensorflow-without-a-phd>]



# Some Improvement

- ▶ Better activation function optimizer, e.g., using  $\text{ReLU}(z) = \max(0, z)$ .
- ▶ Overcome Network overfitting, e.g., using dropout.
- ▶ Network initialization. e.g., using He initialization.
- ▶ Better better optimizer, e.g., using Adam optimizer.

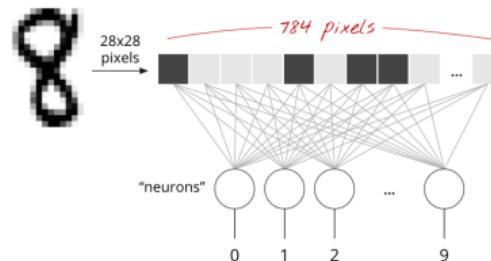


[<https://github.com/GoogleCloudPlatform/tensorflow-without-a-phd>]



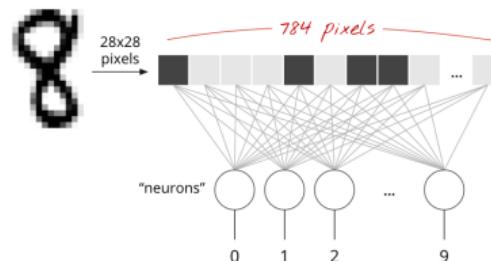
# Vanilla Deep Neural Networks Challenges (1/2)

- ▶ Pixels of each image was flattened into a single vector (really **bad idea**).



# Vanilla Deep Neural Networks Challenges (1/2)

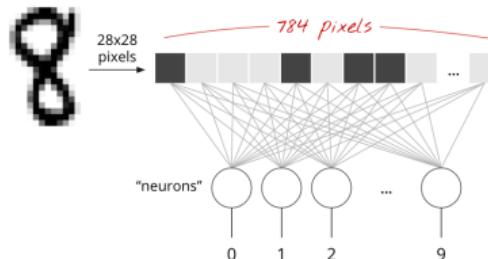
- ▶ Pixels of each image was flattened into a single vector (really **bad idea**).



- ▶ Vanilla deep neural networks **do not scale**.
  - In MNIST, images are black-and-white **28x28** pixel images:  $28 \times 28 = 784$  weights.

# Vanilla Deep Neural Networks Challenges (1/2)

- ▶ Pixels of each image was flattened into a single vector (really **bad idea**).



- ▶ Vanilla deep neural networks **do not scale**.
  - In MNIST, images are black-and-white **28x28** pixel images:  $28 \times 28 = 784$  weights.
- ▶ Handwritten digits are **made of shapes** and we **discarded the shape** information when we flattened the pixels.

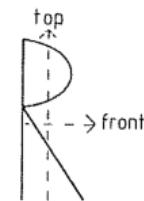
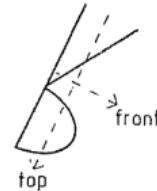


## Vanilla Deep Neural Networks Challenges (2/2)

- ▶ Difficult to recognize objects.

## Vanilla Deep Neural Networks Challenges (2/2)

- ▶ Difficult to **recognize** objects.
- ▶ **Rotation**
- ▶ **Lighting**: objects may **look different** depending on the level of **external lighting**.
- ▶ **Deformation**: objects can be deformed in a variety of **non-affine ways**.
- ▶ **Scale variation**: visual classes often exhibit **variation** in their size.
- ▶ **Viewpoint invariance**.





# Tackle the Challenges

- ▶ Convolutional neural networks (CNN) can tackle the vanilla model challenges.
- ▶ CNN is a type of neural network that can take advantage of shape information.



# Tackle the Challenges

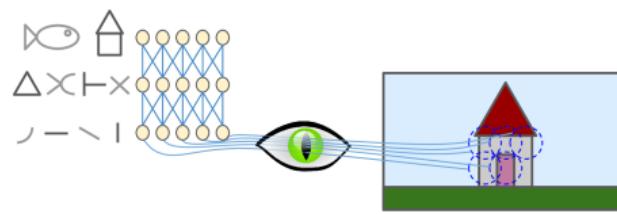
- ▶ Convolutional neural networks (CNN) can tackle the vanilla model challenges.
- ▶ CNN is a type of neural network that can take advantage of shape information.
- ▶ It applies a series of filters to the raw pixel data of an image to extract and learn higher-level features, which the model can then use for classification.



# Filters and Convolution Operations

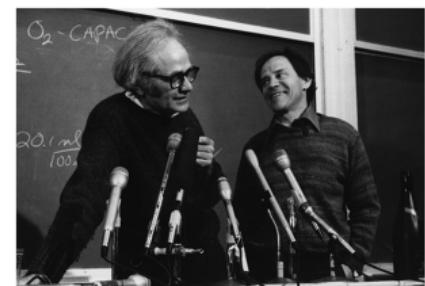
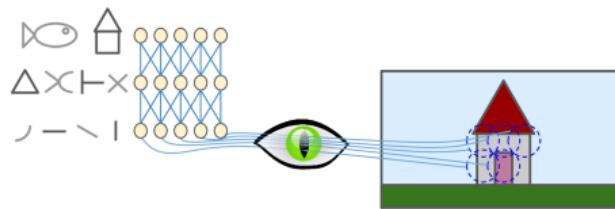
# Brain Visual Cortex Inspired CNNs

- ▶ In 1959, David H. Hubel and Torsten Wiesel.
- ▶ Many **neurons in the visual cortex** have a **small local receptive field**.



# Brain Visual Cortex Inspired CNNs

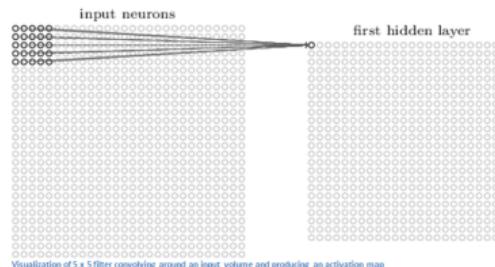
- ▶ In 1959, David H. Hubel and Torsten Wiesel.
- ▶ Many **neurons in the visual cortex** have a **small local receptive field**.
- ▶ They **react** only to visual stimuli located in a **limited region of the visual field**.





# Receptive Fields and Filters

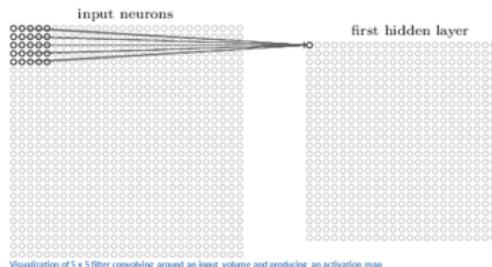
- ▶ Imagine a **flashlight** that is shining over the top left of the image.



[<https://adshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks>]

# Receptive Fields and Filters

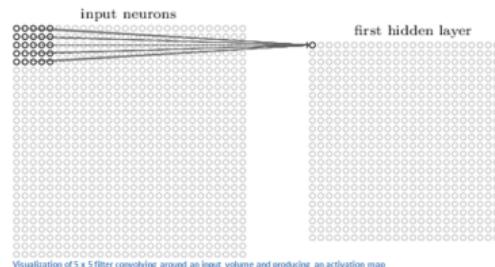
- ▶ Imagine a **flashlight** that is shining over the top left of the image.
- ▶ The **region that it is shining over** is called the **receptive field**.
- ▶ This **flashlight** is called a **filter**.



[<https://adshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks>]

# Receptive Fields and Filters

- ▶ Imagine a **flashlight** that is shining over the top left of the image.
- ▶ The **region that it is shining over** is called the **receptive field**.
- ▶ This **flashlight** is called a **filter**.
- ▶ A filter is a **set of weights**.
- ▶ A **filter** is a **feature detector**, e.g., straight edges, simple colors, and curves.

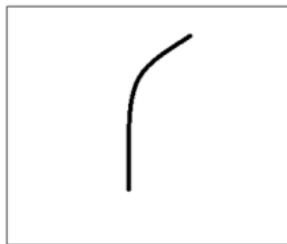


[<https://adshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks>]

# Filters Example (1/3)

0	0	0	0	0	30	0	0
0	0	0	0	30	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	0	0	0	0	0

Pixel representation of filter



Visualization of a curve detector filter

## Filters Example (1/3)

0	0	0	0	0	30	0	0
0	0	0	0	30	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	30	0	0	0	0
0	0	0	0	0	0	0	0

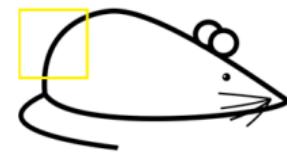
Pixel representation of filter



Visualization of a curve detector filter



Original image



Visualization of the filter on the image

[<https://adेशपांडे3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks>]

## Filters Example (2/3)



Visualization of the receptive field

0	0	0	0	0	0	30	0
0	0	0	0	50	50	50	0
0	0	0	20	50	0	0	0
0	0	0	50	50	0	0	0
0	0	0	50	50	0	0	0
0	0	0	50	50	0	0	0
0	0	0	50	50	0	0	0
0	0	0	50	50	0	0	0

Pixel representation of the receptive field

\*

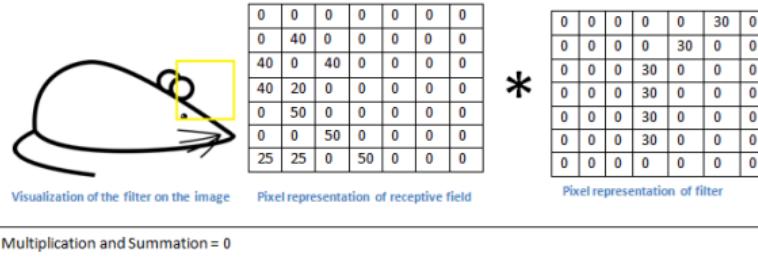
0	0	0	0	0	30	0
0	0	0	0	30	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	30	0	0	0
0	0	0	0	0	0	0

Pixel representation of filter

$$\text{Multiplication and Summation} = (50 \cdot 30) + (50 \cdot 30) + (50 \cdot 30) + (20 \cdot 30) + (50 \cdot 30) = 6600 \text{ (A large number!)}$$

[<https://adephande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks>]

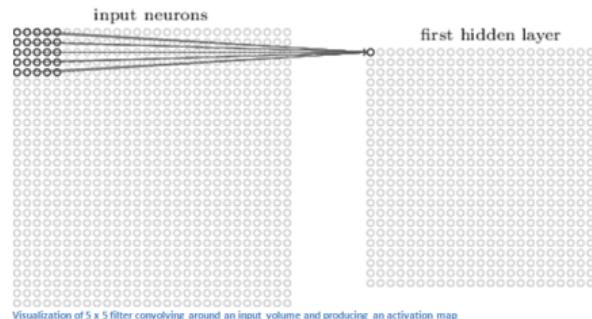
## Filters Example (3/3)



[<https://adephante3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks>]

# Convolution Operation

- ▶ Convolution takes a **filter** and multiplying it over the entire area of an input image.
- ▶ Imagine this **flashlight (filter)** sliding across all the areas of the input image.



[<https://adephpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks>]



## Convolution Operation - More Formal Definition

- ▶ Convolution is a mathematical operation on two functions  $x$  and  $h$ .
  - You can think of  $x$  as the input image, and  $h$  as a filter (kernel) on the input image.



## Convolution Operation - More Formal Definition

- ▶ Convolution is a mathematical operation on two functions  $x$  and  $h$ .
  - You can think of  $x$  as the input image, and  $h$  as a filter (kernel) on the input image.
- ▶ For a 1D convolution we can define it as below:

$$y(k) = \sum_{n=0}^{N-1} h(n) \cdot x(k-n)$$

- ▶  $N$  is the number of elements in  $h$ .



## Convolution Operation - More Formal Definition

- ▶ Convolution is a mathematical operation on two functions  $x$  and  $h$ .
  - You can think of  $x$  as the input image, and  $h$  as a filter (kernel) on the input image.
- ▶ For a 1D convolution we can define it as below:

$$y(k) = \sum_{n=0}^{N-1} h(n) \cdot x(k-n)$$

- ▶  $N$  is the number of elements in  $h$ .
- ▶ We are sliding the filter  $h$  over the input image  $x$ .



## Convolution Operation - 1D Example (1/2)

- ▶ Suppose our input 1D image is  $x$ , and filter  $h$  are as follows:

$$x = \boxed{10 \quad 50 \quad 60 \quad 10 \quad 20 \quad 40 \quad 30}$$

$$h = \boxed{1/3 \quad 1/3 \quad 1/3}$$

- ▶ Let's call the output image  $y$ .
- ▶ What is the value of  $y(3)$ ?

$$y(k) = \sum_{n=0}^{N-1} h(n) \cdot x(k-n)$$

## Convolution Operation - 1D Example (2/2)

- To compute  $y(3)$ , we slide the filter so that it is centered around  $x(3)$ .

10	50	60	10	20	30	40
0	1/3	1/3	1/3	0	0	0

$$y(3) = \frac{1}{3}50 + \frac{1}{3}60 + \frac{1}{3}10 = 40$$

## Convolution Operation - 1D Example (2/2)

- To compute  $y(3)$ , we slide the filter so that it is centered around  $x(3)$ .

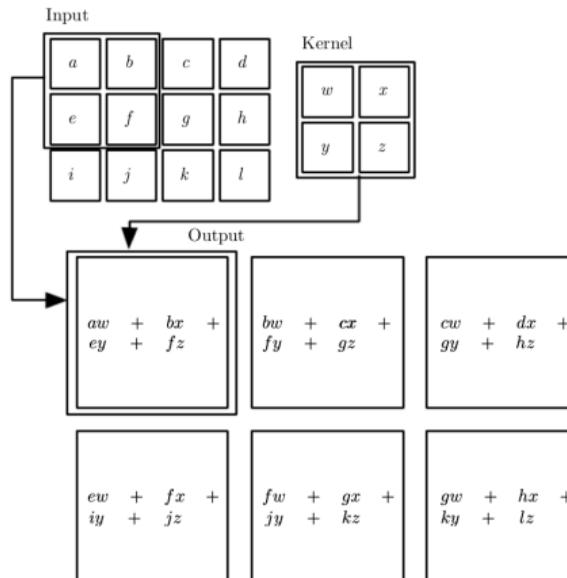
10	50	60	10	20	30	40
0	1/3	1/3	1/3	0	0	0

$$y(3) = \frac{1}{3}50 + \frac{1}{3}60 + \frac{1}{3}10 = 40$$

- We can compute the other values of  $y$  as well.

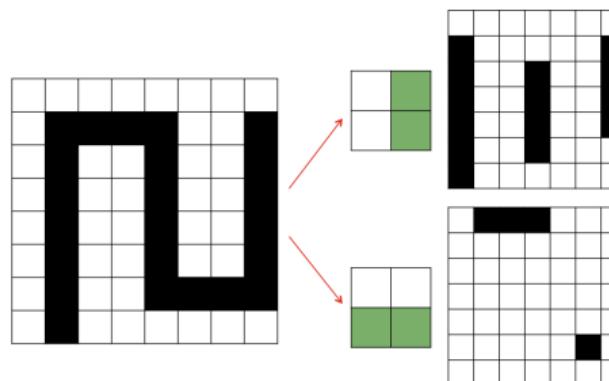
$$y = [20 \ 40 \ 40 \ 30 \ 20 \ 30 \ 23.333]$$

# Convolution Operation - 2D Example (1/2)



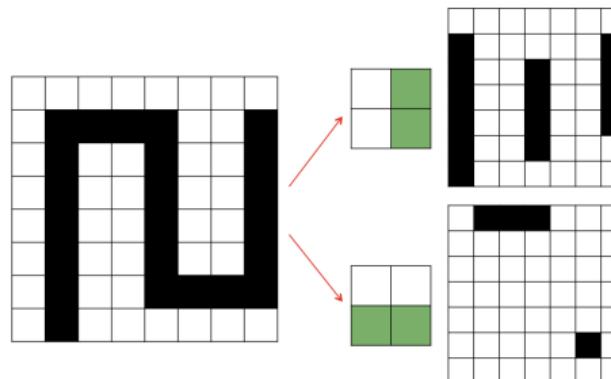
## Convolution Operation - 2D Example (2/2)

- ▶ Detect **vertical** and **horizontal lines** in an image.
- ▶ **Slide the filters** across the entirety of the image.



## Convolution Operation - 2D Example (2/2)

- ▶ Detect **vertical** and **horizontal lines** in an image.
- ▶ **Slide the filters** across the entirety of the image.
- ▶ The **result** is our **feature map**: indicates where we've found the **feature** we're looking for in the original image.

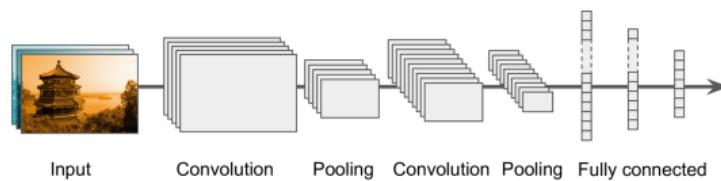




# Convolutional Neural Network (CNN)

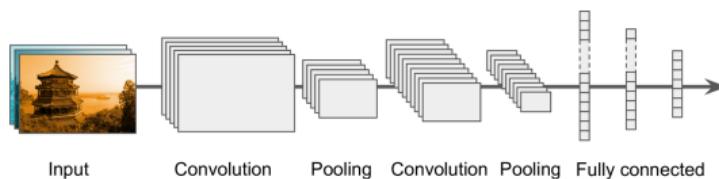
# CNN Components (1/2)

- ▶ **Convolutional layers**: apply a specified number of **convolution filters** to the image.



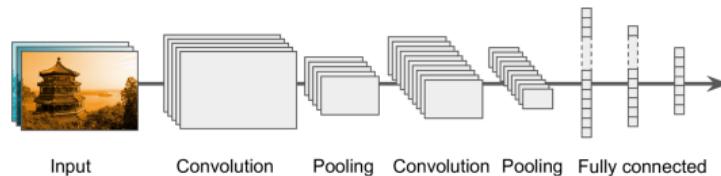
# CNN Components (1/2)

- ▶ **Convolutional layers**: apply a specified number of **convolution filters** to the image.
- ▶ **Pooling layers**: **downsample the image** data extracted by the convolutional layers to **reduce the dimensionality** of the feature map in order to decrease processing time.



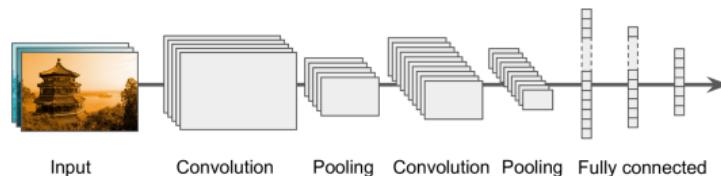
## CNN Components (1/2)

- ▶ **Convolutional layers**: apply a specified number of **convolution filters** to the image.
- ▶ **Pooling layers**: **downsample the image** data extracted by the convolutional layers to **reduce the dimensionality** of the feature map in order to decrease processing time.
- ▶ **Dense layers**: a **fully connected layer** that performs **classification** on the features extracted by the convolutional layers and downsampled by the pooling layers.



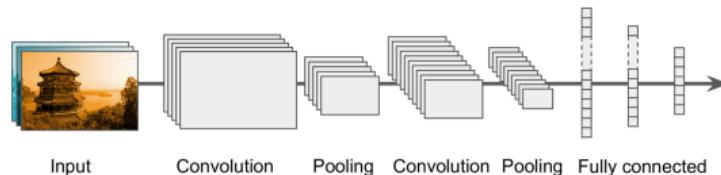
## CNN Components (2/2)

- ▶ A CNN is composed of a stack of convolutional modules.



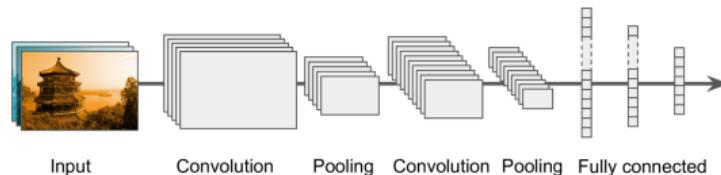
## CNN Components (2/2)

- ▶ A **CNN** is composed of a **stack of convolutional modules**.
- ▶ Each **module** consists of a **convolutional layer** followed by a **pooling layer**.



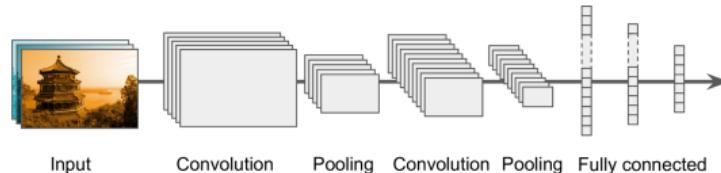
## CNN Components (2/2)

- ▶ A **CNN** is composed of a **stack of convolutional modules**.
- ▶ Each **module** consists of a **convolutional layer** followed by a **pooling layer**.
- ▶ The **last module** is followed by **one or more dense layers** that perform **classification**.

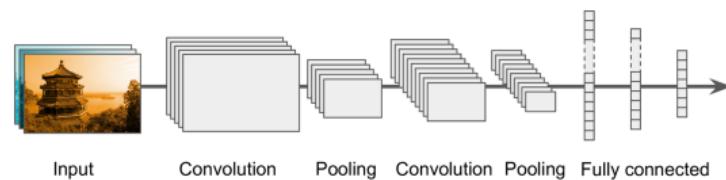


## CNN Components (2/2)

- ▶ A **CNN** is composed of a **stack of convolutional modules**.
- ▶ Each **module** consists of a **convolutional layer** followed by a **pooling layer**.
- ▶ The **last module** is followed by **one or more dense layers** that perform **classification**.
- ▶ The **final dense layer** contains a **single node** for each target class in the model, with a **softmax** activation function.

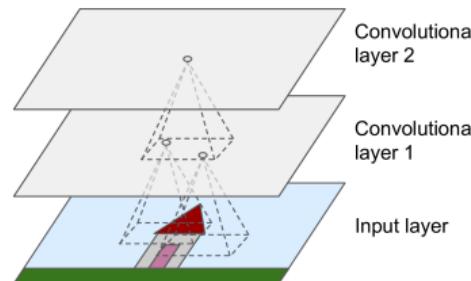


# Convolutional Layer



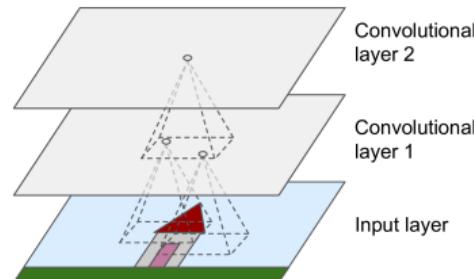
# Convolutional Layer (1/4)

- ▶ Sparse interactions
- ▶ Each neuron in the convolutional layers are **only** connected to pixels in their **receptive fields** (not every single pixel).



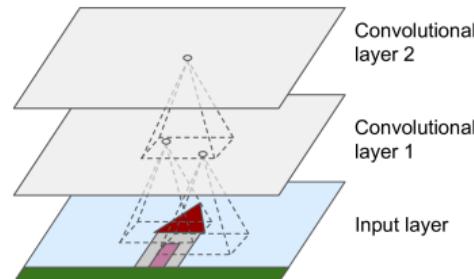
## Convolutional Layer (2/4)

- ▶ Each neuron applies **filters** on its **receptive field**.



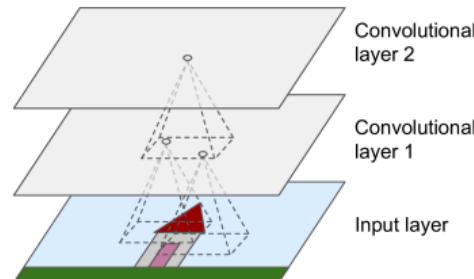
## Convolutional Layer (2/4)

- ▶ Each neuron applies **filters** on its **receptive field**.
  - Calculates a **weighted sum** of the input pixels in the receptive fields



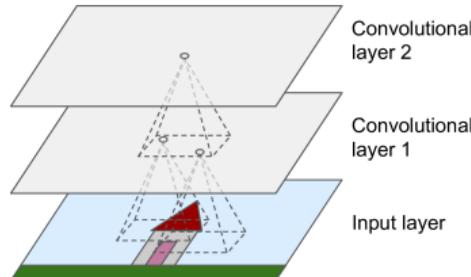
## Convolutional Layer (2/4)

- ▶ Each neuron applies **filters** on its **receptive field**.
  - Calculates a **weighted sum** of the input pixels in the receptive fields
- ▶ Adds a **bias**, and feeds the result through its **activation function** to the next layer.



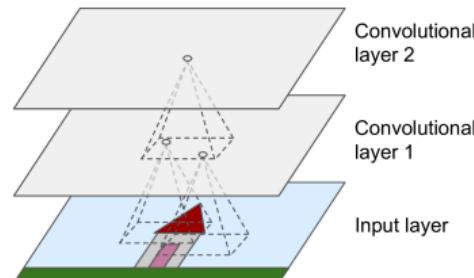
## Convolutional Layer (2/4)

- ▶ Each neuron applies **filters** on its **receptive field**.
  - Calculates a **weighted sum** of the input pixels in the receptive fields
- ▶ Adds a **bias**, and feeds the result through its **activation function** to the next layer.
- ▶ The **output** of this layer is a **feature map (activation map)**



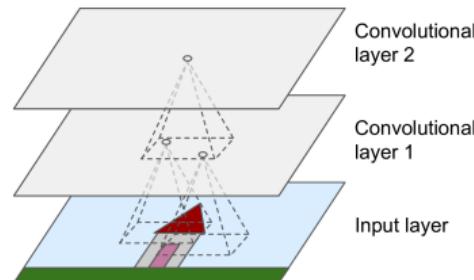
## Convolutional Layer (3/4)

- ▶ Parameter sharing
- ▶ All neurons of a convolutional layer reuses the same weights.



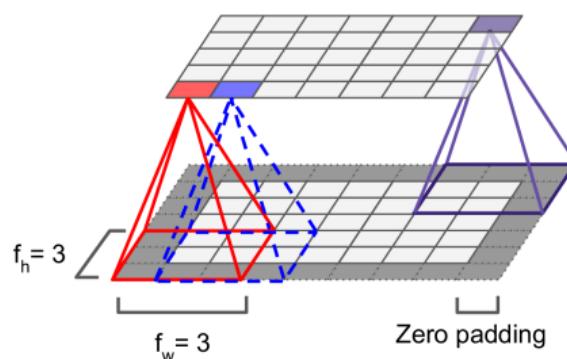
## Convolutional Layer (3/4)

- ▶ Parameter sharing
- ▶ All neurons of a convolutional layer reuses the same weights.
- ▶ They apply the same filter in different positions.
- ▶ Whereas in a fully-connected network, each neuron had its own set of weights.



## Convolutional Layer (4/4)

- ▶ Assume the filter size (kernel size) is  $f_w \times f_h$ .
  - $f_h$  and  $f_w$  are the height and width of the receptive field, respectively.
- ▶ A neuron in row  $i$  and column  $j$  of a given layer is connected to the outputs of the neurons in the previous layer in rows  $i$  to  $i + f_h - 1$ , and columns  $j$  to  $j + f_w - 1$ .



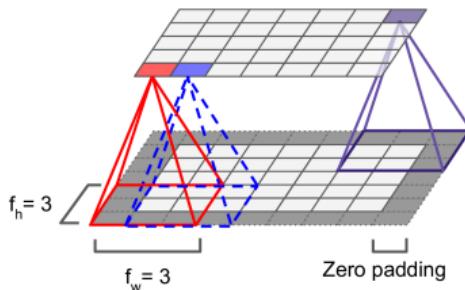


# Padding

- ▶ What will happen if you apply a **5x5 filters** to a **32x32 input** volume?
  - The output volume would be **28x28**.
  - The spatial **dimensions decrease**.

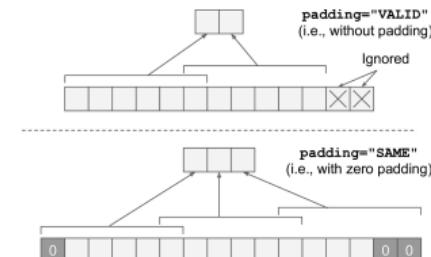
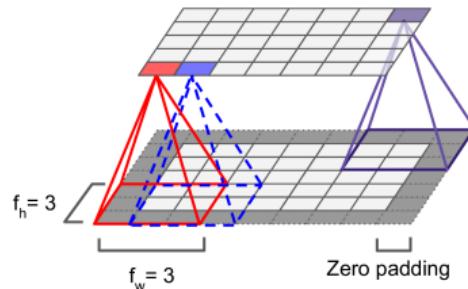
# Padding

- ▶ What will happen if you apply a **5x5 filters** to a **32x32 input volume**?
  - The output volume would be **28x28**.
  - The spatial **dimensions decrease**.
- ▶ **Zero padding:** in order for a layer to have the **same height and width** as the previous layer, it is common to **add zeros around the inputs**.



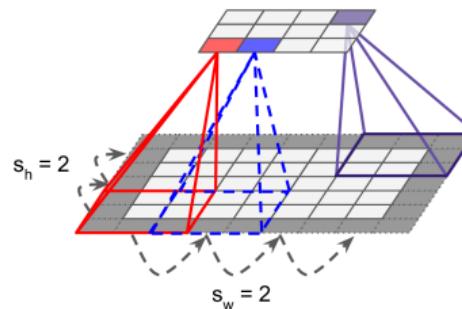
# Padding

- ▶ What will happen if you apply a **5x5 filters** to a **32x32 input volume**?
  - The output volume would be **28x28**.
  - The spatial **dimensions decrease**.
- ▶ **Zero padding**: in order for a layer to have the **same height and width** as the previous layer, it is common to **add zeros around the inputs**.
- ▶ In **Tensorflow**, padding can be either **SAME** or **VALID** to have zero padding or not.



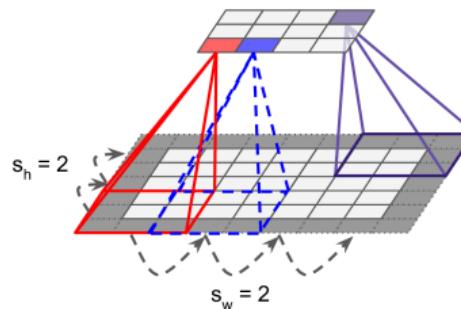
## Stride (1/2)

- ▶ The **distance** between two consecutive receptive fields is called the **stride**.



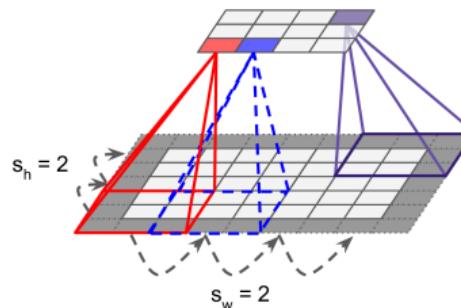
## Stride (1/2)

- ▶ The **distance** between two consecutive receptive fields is called the **stride**.
- ▶ The stride controls **how the filter convolves** around the input volume.

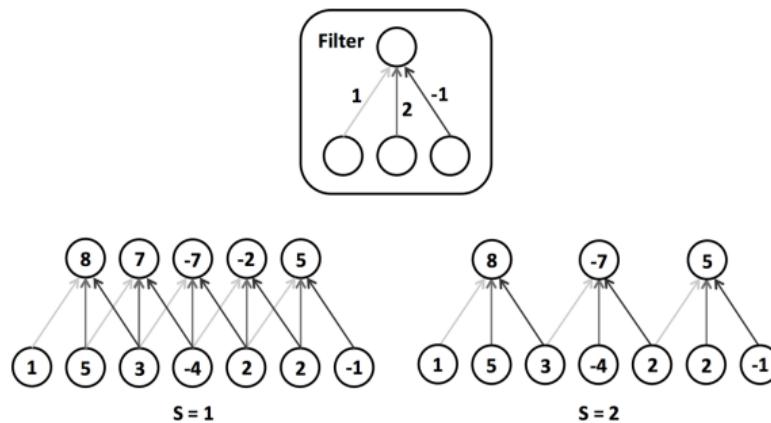


## Stride (1/2)

- ▶ The **distance** between two consecutive receptive fields is called the **stride**.
- ▶ The stride controls **how the filter convolves** around the input volume.
- ▶ Assume  $s_h$  and  $s_w$  are the **vertical and horizontal strides**, then, a neuron located in **row  $i$**  and **column  $j$**  in a layer is connected to the outputs of the neurons in the **previous layer** located in **rows  $i \times s_h$**  to  **$i \times s_h + f_h - 1$** , and **columns  $j \times s_w$**  to  **$j \times s_w + f_w - 1$** .

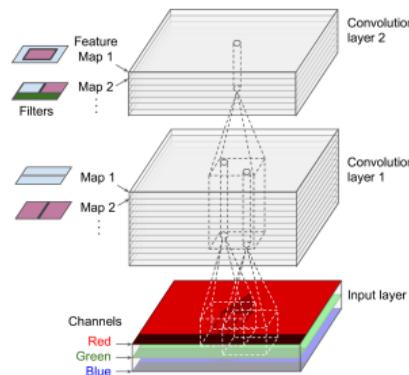


## Stride (2/2)



# Stacking Multiple Feature Maps

- ▶ Up to now, we represented each convolutional layer with a **single feature map**.
- ▶ Each convolutional layer can be composed of **several feature maps** of equal sizes.
- ▶ Input images are also composed of **multiple sublayers**: **one per color channel**.
- ▶ A **convolutional layer simultaneously applies multiple filters** to its inputs.

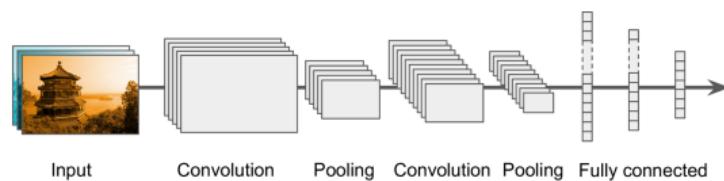




## Activation Function

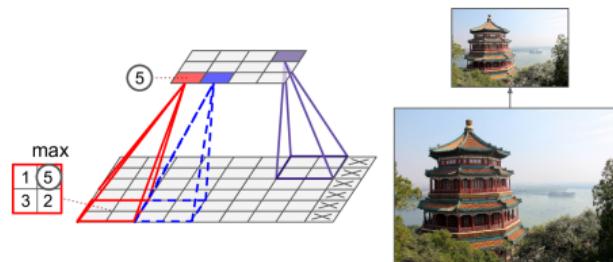
- ▶ After calculating a **weighted sum** of the input pixels in the **receptive fields**, and adding **biases**, each neuron feeds the result through its **ReLU activation function** to the next layer.
- ▶ The purpose of this activation function is to add **non-linearity** to a system.

# Pooling Layer



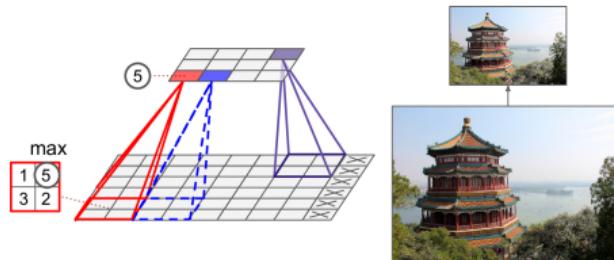
## Pooling Layer (1/2)

- ▶ After the activation functions, we can apply a **pooling layer**.
- ▶ Its goal is to **subsample (shrink)** the input image.



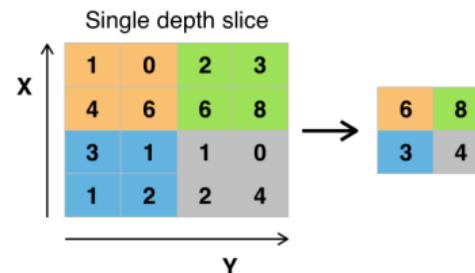
## Pooling Layer (1/2)

- ▶ After the activation functions, we can apply a **pooling layer**.
- ▶ Its goal is to **subsample (shrink)** the input image.
  - To **reduce** the computational load, the memory usage, and the number of parameters.



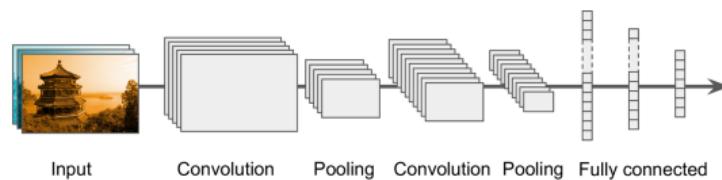
## Pooling Layer (2/2)

- ▶ Each **neuron** in a pooling layer is connected to the outputs of a **receptive field** in the previous layer.
- ▶ A pooling neuron has **no weights**.
- ▶ It **aggregates** the inputs using an aggregation function such as the **max or mean**.



Example of Maxpool with a 2x2 filter and a stride of 2

# Fully Connected Layer





## Fully Connected Layer

- ▶ This layer takes an input from the **last convolution module**, and outputs an **N** dimensional vector.
  - **N** is the **number of classes** that the program has to choose from.



## Fully Connected Layer

- ▶ This layer takes an input from the **last convolution module**, and outputs an **N** dimensional vector.
  - **N** is the **number of classes** that the program has to choose from.
- ▶ For example, if you wanted a **digit classification** program, **N** would be 10.



## Fully Connected Layer

- ▶ This layer takes an input from the **last convolution module**, and outputs an **N** dimensional vector.
  - **N** is the **number of classes** that the program has to choose from.
- ▶ For example, if you wanted a **digit classification** program, **N** would be 10.
- ▶ Each number in this **N** dimensional vector represents the **probability of a certain class**.



# Flattening

- ▶ We need to **convert the output** of the convolutional part of the CNN into a **1D feature vector**.
- ▶ This operation is called **flattening**.



# Flattening

- ▶ We need to **convert the output** of the convolutional part of the CNN into a **1D feature vector**.
- ▶ This operation is called **flattening**.
- ▶ It gets the **output of the convolutional layers**, **flattens** all its structure to create a **single long feature vector** to be used by the **dense layer** for the final classification.



# Example

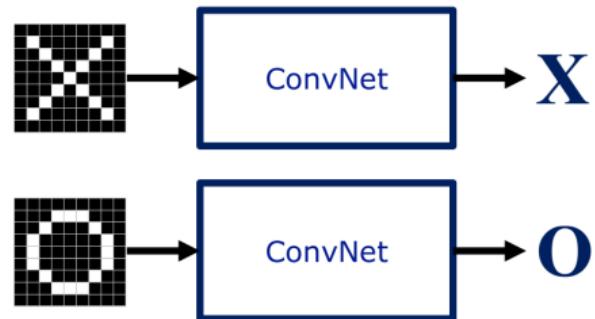


## A Toy ConvNet: X's and O's

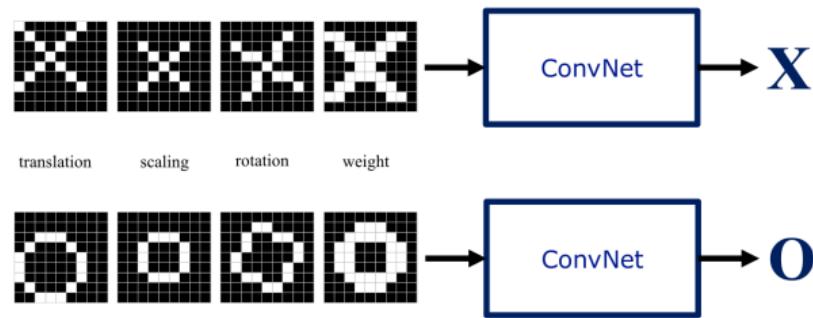
A two-dimensional  
array of pixels



## For Example

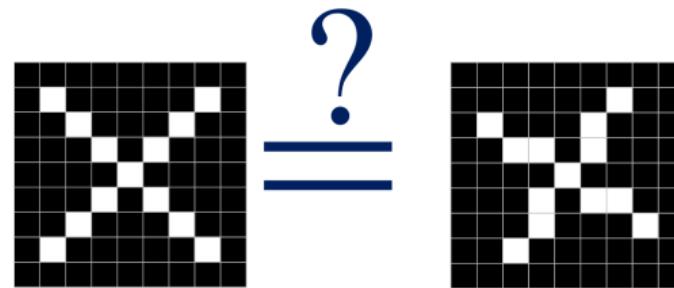


## Trickier Cases





# Deciding is Hard





# What Computers See



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	1	-1
-1	1	1	-1	-1	-1	1	-1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	-1	-1	1	-1	-1	-1	-1	
-1	-1	-1	1	-1	1	-1	-1	-1	
-1	-1	1	-1	-1	-1	1	-1	-1	
-1	1	-1	-1	-1	-1	1	-1	-1	
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

# What Computers See



-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	1	-1	-1
-1	-1	-1	1	-1	1	1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	X	-1	-1	-1	-1	X	X	-1
-1	X	X	-1	-1	X	X	-1	-1
-1	-1	X	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	X	-1	-1
-1	-1	X	X	-1	-1	X	X	-1
-1	X	X	-1	-1	-1	X	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

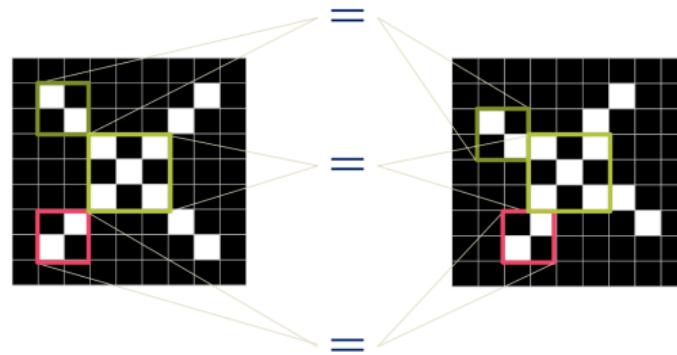
# Computers are Literal

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	1



-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	1	-1

# ConvNets Match Pieces of the Image



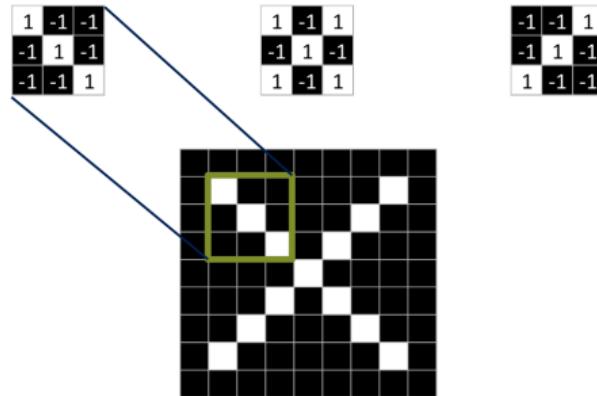
## Filters Match Pieces of the Image

$$\begin{array}{|c|c|c|} \hline 1 & -1 & -1 \\ \hline -1 & 1 & -1 \\ \hline -1 & -1 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline 1 & -1 & 1 \\ \hline -1 & 1 & -1 \\ \hline 1 & -1 & 1 \\ \hline \end{array}$$

$$\begin{array}{|c|c|c|} \hline -1 & -1 & 1 \\ \hline -1 & 1 & -1 \\ \hline 1 & -1 & -1 \\ \hline \end{array}$$

# Filters Match Pieces of the Image

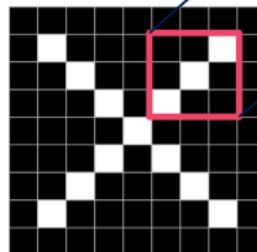


# Filters Match Pieces of the Image

$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} -1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & -1 \end{bmatrix}$$

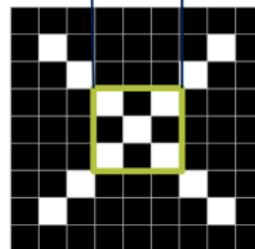


# Filters Match Pieces of the Image

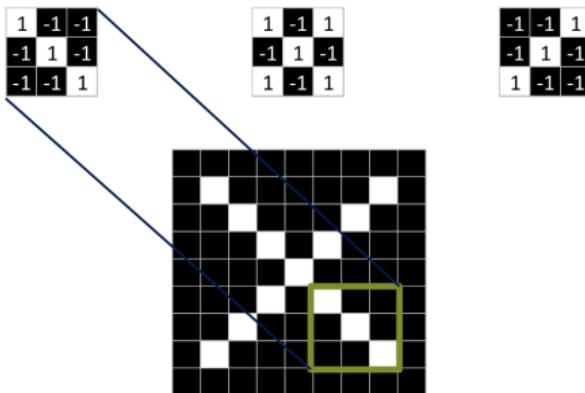
$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} -1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & -1 \end{bmatrix}$$



# Filters Match Pieces of the Image

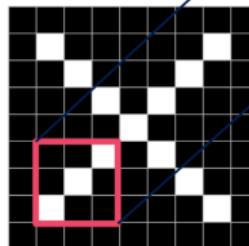


# Filters Match Pieces of the Image

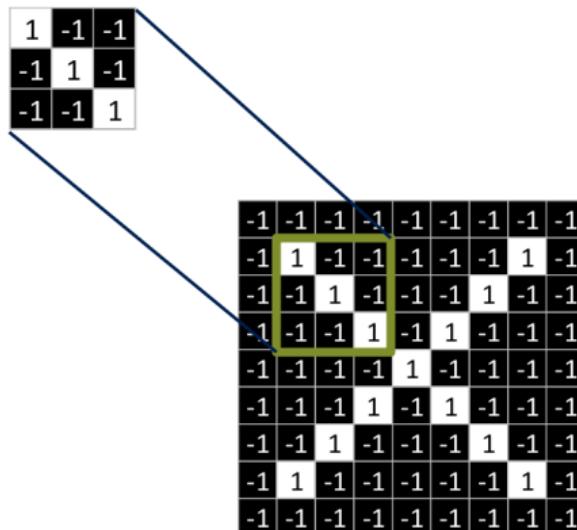
$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & 1 \end{bmatrix}$$

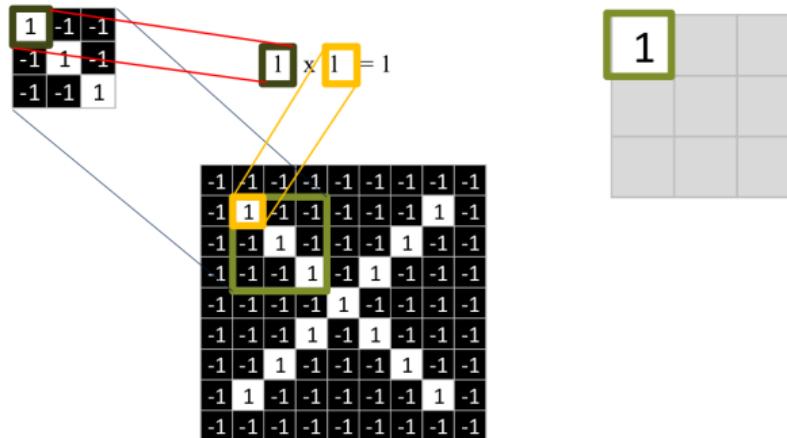
$$\begin{bmatrix} -1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & -1 \end{bmatrix}$$



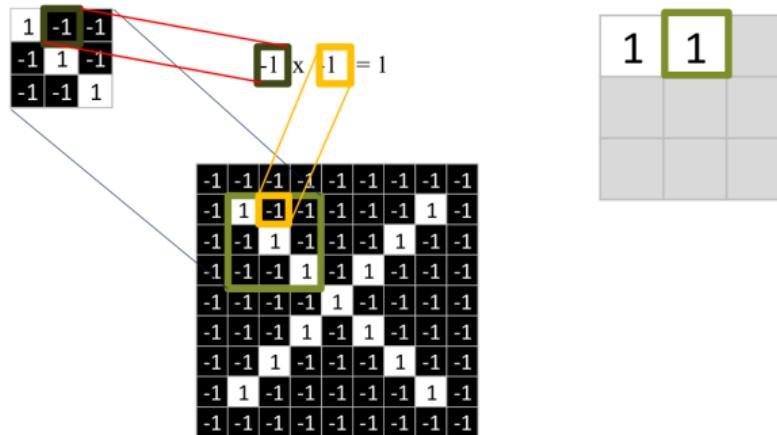
# Filtering: The Math Behind the Match



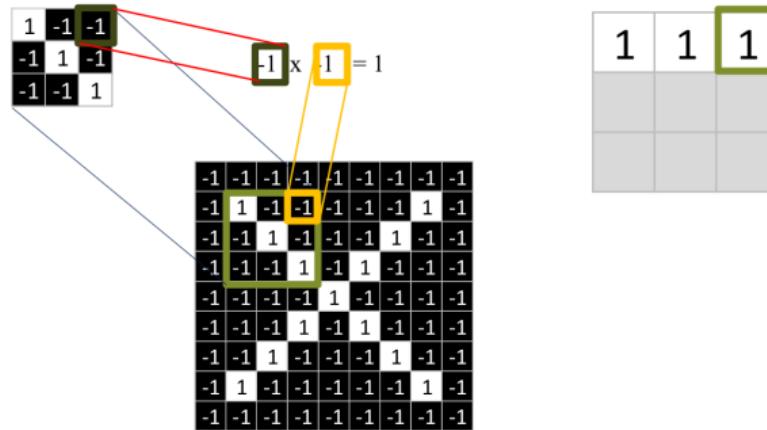
# Filtering: The Math Behind the Match



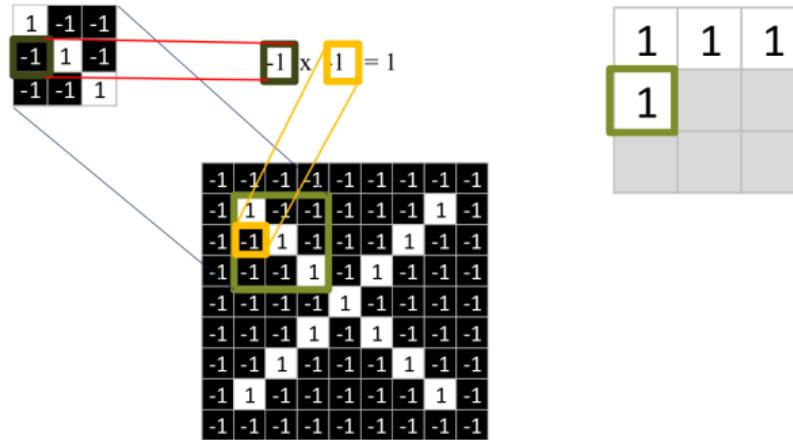
# Filtering: The Math Behind the Match



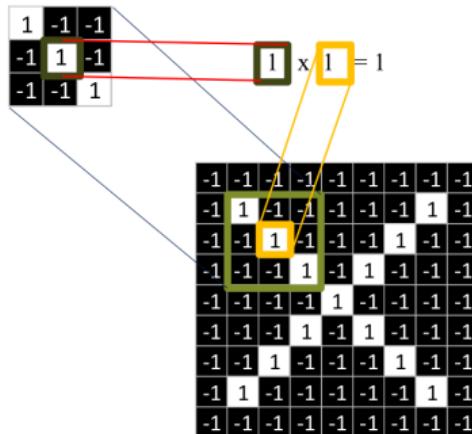
# Filtering: The Math Behind the Match



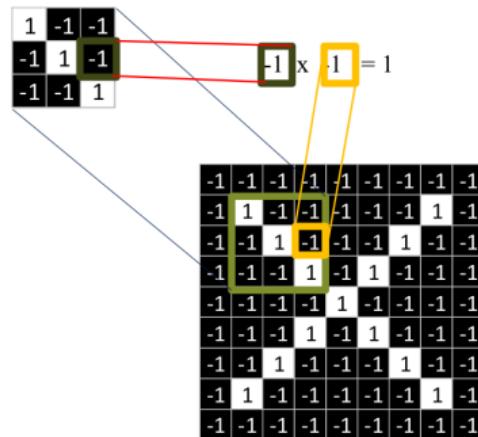
# Filtering: The Math Behind the Match



# Filtering: The Math Behind the Match

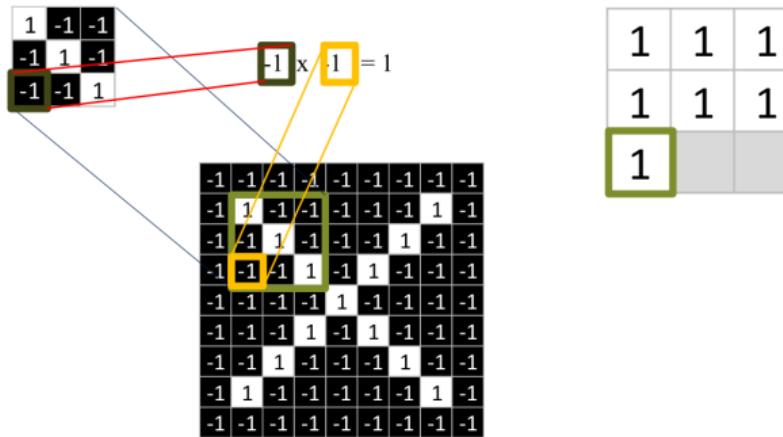


# Filtering: The Math Behind the Match

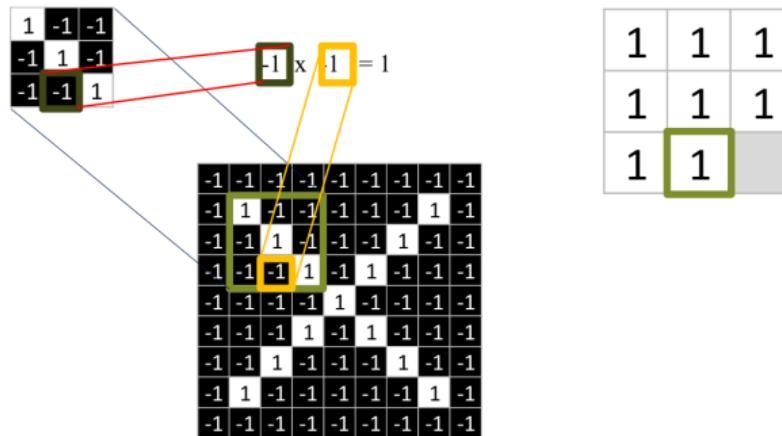


1	1	1
1	1	1
1	1	1

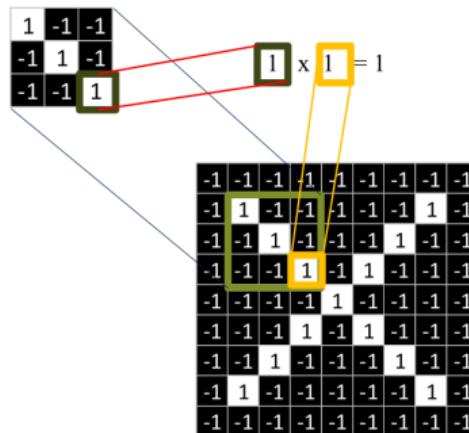
# Filtering: The Math Behind the Match



# Filtering: The Math Behind the Match

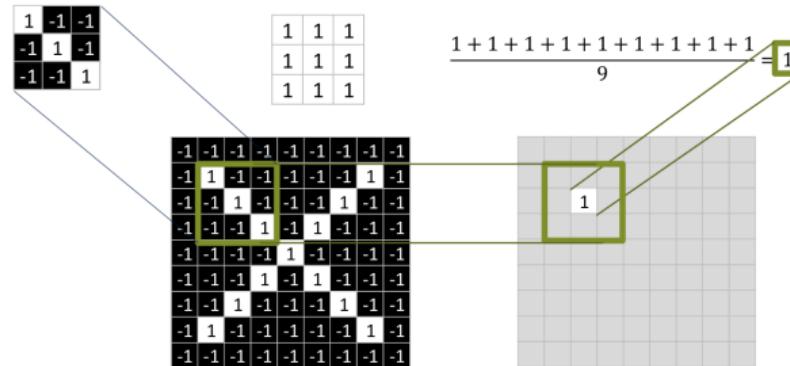


# Filtering: The Math Behind the Match

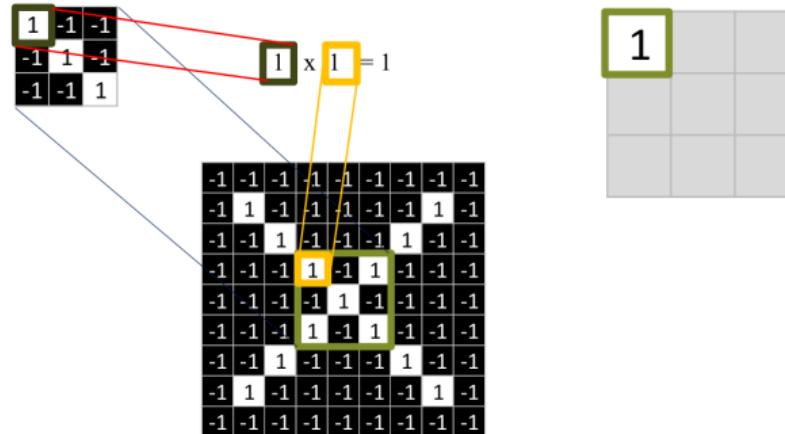


1	1	1
1	1	1
1	1	1

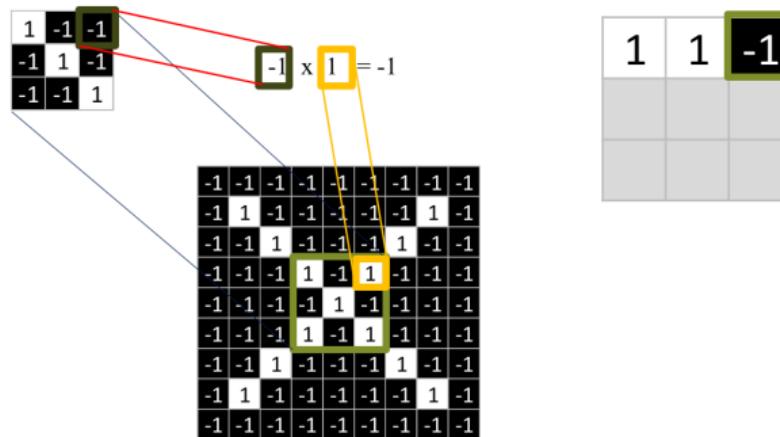
# Filtering: The Math Behind the Match



# Filtering: The Math Behind the Match

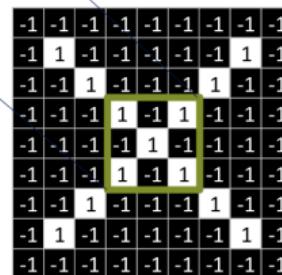


# Filtering: The Math Behind the Match



# Filtering: The Math Behind the Match

$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix}$$

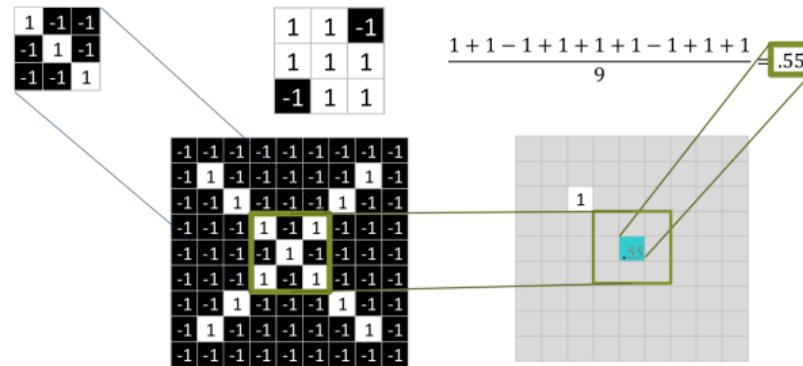


A diagram illustrating the application of a convolutional filter to an input image. A small 3x3 kernel matrix is shown at the top left. An arrow points from this kernel to a larger 13x13 input matrix below it. The input matrix has values ranging from -1 to 1. A 3x3 submatrix in the center of the input is highlighted with a green border, representing the receptive field of a single output unit. The result of the convolution is shown on the right, where the output value is 1, indicated by a black box around the corresponding entry in the result matrix.

$$\begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & 1 & -1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & 1 & 1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 & -1 \\ -1 & -1 & -1 & -1 & -1 & -1 & 1 & -1 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & -1 \\ 1 & 1 & 1 \\ -1 & 1 & 1 \end{bmatrix}$$

# Filtering: The Math Behind the Match



# Convolution: Trying Every Possible Match

$$\begin{array}{|c|c|c|c|c|c|c|c|c|} \hline -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ \hline -1 & 1 & -1 & -1 & -1 & -1 & -1 & 1 & -1 \\ \hline -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 & -1 \\ \hline -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 & -1 \\ \hline -1 & -1 & -1 & -1 & 1 & -1 & -1 & -1 & -1 \\ \hline -1 & -1 & -1 & -1 & 1 & -1 & -1 & -1 & -1 \\ \hline -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 & -1 \\ \hline -1 & -1 & 1 & -1 & -1 & 1 & -1 & -1 & -1 \\ \hline -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 & -1 \\ \hline -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & -1 & -1 \\ \hline -1 & 1 & -1 \\ \hline -1 & -1 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|c|c|c|c|c|c|} \hline 0.77 & -0.11 & 0.11 & 0.33 & 0.55 & -0.11 & 0.33 \\ \hline -0.11 & 1.00 & -0.11 & 0.33 & -0.11 & 0.11 & -0.11 \\ \hline 0.11 & -0.11 & 1.00 & -0.33 & 0.11 & -0.11 & 0.55 \\ \hline 0.33 & 0.33 & -0.33 & 0.55 & -0.33 & 0.33 & 0.33 \\ \hline 0.55 & -0.11 & 0.11 & -0.33 & 1.00 & -0.11 & 0.11 \\ \hline -0.11 & 0.11 & -0.11 & 0.33 & -0.11 & 1.00 & -0.11 \\ \hline 0.33 & -0.11 & 0.55 & 0.33 & 0.11 & -0.11 & 0.77 \\ \hline \end{array}$$

# Three Filters Here, So Three Images Out

$$\begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & 1 & -1 & -1 & 1 & 1 & -1 \\ -1 & -1 & 1 & 1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & -1 & -1 \\ -1 & -1 & 1 & 1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 \\ -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & -1 & -1 & -1 & -1 \end{bmatrix}$$



$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix}$$

=

$$\begin{bmatrix} 0.77 & -0.11 & 0.11 & 0.33 & 0.55 & -0.11 & 0.33 \\ -0.11 & 1.00 & -0.11 & 0.33 & -0.11 & 0.11 & -0.11 \\ 0.11 & -0.11 & 1.00 & 0.33 & 0.33 & -0.11 & 0.33 \\ 0.33 & 0.11 & -0.33 & 0.55 & -0.33 & 0.33 & 0.33 \\ 0.55 & -0.11 & 0.11 & 0.33 & 1.00 & -0.11 & 0.11 \\ -0.11 & 0.11 & -0.11 & -0.33 & -0.11 & 1.00 & -0.11 \\ 0.33 & -0.11 & 0.55 & 0.33 & 0.33 & -0.11 & 0.77 \end{bmatrix}$$

$$\begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & 1 & -1 & 1 & 1 & -1 \\ -1 & -1 & 1 & -1 & 1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & -1 & -1 \\ -1 & -1 & 1 & 1 & 1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 \\ -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 \end{bmatrix}$$



$$\begin{bmatrix} 1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & 1 \end{bmatrix}$$

=

$$\begin{bmatrix} 0.33 & -0.55 & 0.11 & 0.11 & -0.55 & 0.33 \\ -0.55 & 0.55 & -0.55 & 0.33 & -0.55 & 0.55 & -0.55 \\ 0.11 & -0.55 & 0.55 & 0.77 & 0.55 & -0.55 & 0.11 \\ -0.11 & 0.33 & -0.77 & 1.00 & 0.77 & 0.33 & -0.33 \\ 0.11 & -0.55 & 0.55 & -0.77 & 0.55 & -0.55 & 0.11 \\ -0.55 & 0.55 & -0.55 & 0.11 & -0.55 & 0.55 & -0.55 \\ 0.33 & -0.55 & 0.11 & 0.11 & -0.55 & 0.33 \end{bmatrix}$$

$$\begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & 1 & -1 & 1 & 1 & -1 \\ -1 & -1 & 1 & 1 & 1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & -1 & -1 \\ -1 & -1 & 1 & -1 & 1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 \\ -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 \end{bmatrix}$$



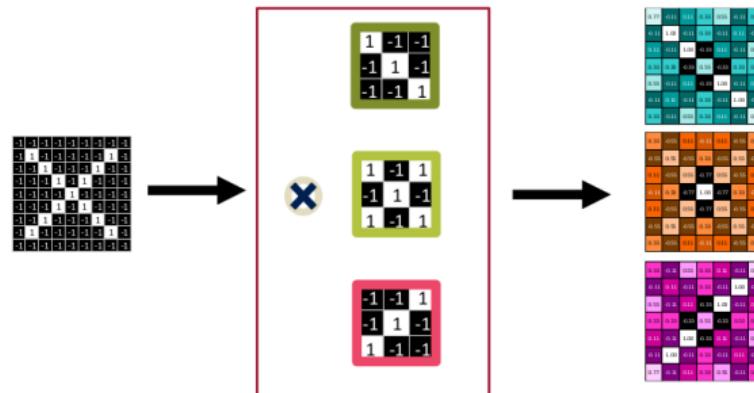
$$\begin{bmatrix} -1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & -1 \end{bmatrix}$$

=

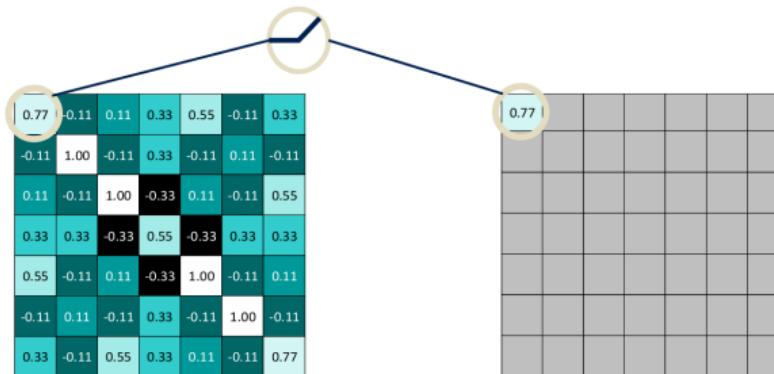
$$\begin{bmatrix} 0.33 & -0.11 & 0.55 & 0.33 & 0.11 & -0.11 & 0.77 \\ -0.11 & 0.11 & -0.11 & 0.33 & 0.11 & 1.00 & 0.11 \\ 0.55 & -0.11 & 0.11 & 0.33 & 1.00 & -0.11 & 0.11 \\ 0.33 & 0.11 & -0.33 & 0.55 & 0.33 & 0.33 & 0.33 \\ 0.11 & -0.11 & 1.00 & 0.33 & 0.11 & -0.11 & 0.55 \\ -0.11 & 1.00 & -0.11 & 0.33 & 0.11 & 0.11 & -0.11 \\ 0.77 & -0.11 & 0.11 & 0.33 & 0.55 & -0.11 & 0.33 \end{bmatrix}$$

# Convolution Layer

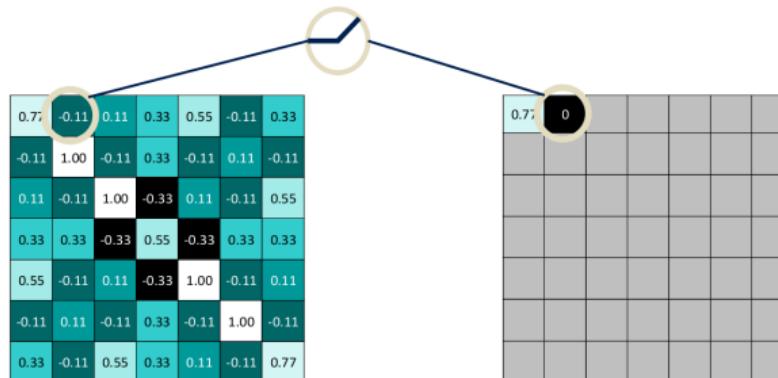
- One image becomes a **stack of filtered images**.



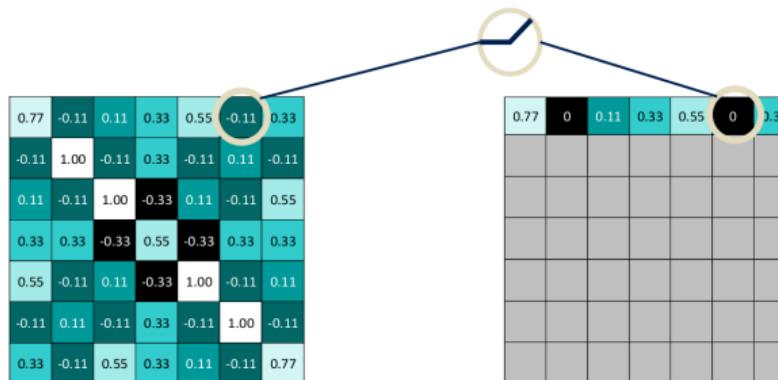
# Rectified Linear Units (ReLUs)



# Rectified Linear Units (ReLUs)



# Rectified Linear Units (ReLUs)



# Rectified Linear Units (ReLUs)

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

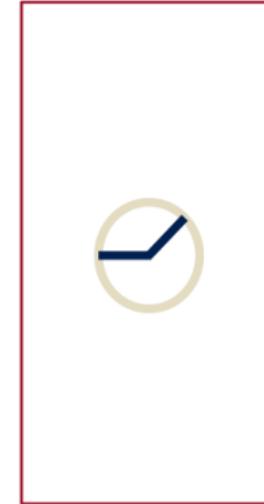


0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77

## ReLU Layer

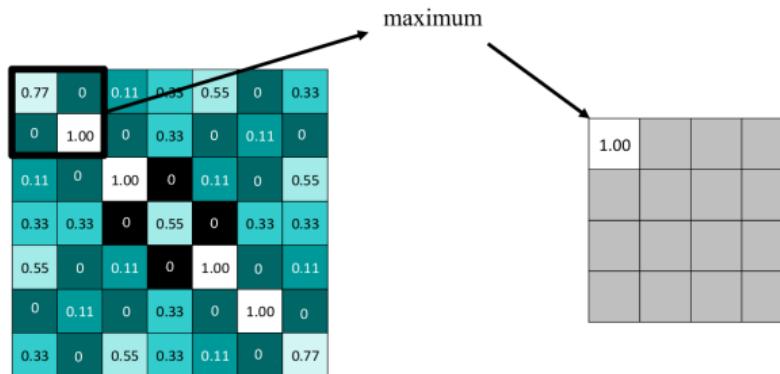
- ▶ A stack of images becomes a stack of images with **no negative values**.

0.77	-0.15	0.11	0.33	-0.05	-0.03	0.31
-0.11	1.09	-0.15	-0.31	-0.11	0.11	-0.11
0.11	-0.11	1.09	0.33	0.11	-0.15	0.35
0.01	0.15	-0.31	-0.51	0.11	0.11	0.31
0.35	-0.11	0.11	0.33	1.09	-0.11	0.31
0.11	0.11	-0.11	0.33	-0.11	1.09	-0.11
0.33	-0.11	0.55	-0.31	0.11	-0.11	0.27
0.33	-0.15	0.11	0.33	0.11	-0.25	0.31
0.36	-0.15	0.55	0.33	-0.15	0.39	0.36
0.33	-0.15	0.55	0.77	0.35	-0.15	0.31
0.11	0.11	-0.77	1.09	0.77	0.11	0.11
0.11	-0.15	0.55	0.77	0.35	-0.15	0.11
0.35	0.55	-0.31	0.33	-0.15	0.35	0.35
0.33	-0.15	0.11	0.33	0.11	-0.25	0.31
0.33	-0.11	0.55	0.33	0.11	-0.11	0.27
0.11	0.11	-0.33	-0.31	0.11	1.09	0.11
0.35	-0.11	0.11	-0.31	1.09	-0.11	0.11
0.33	0.75	-0.33	-0.31	0.33	0.11	0.33
0.33	-0.11	1.09	0.33	0.11	-0.11	0.35
0.11	1.09	0.11	0.33	0.11	0.11	0.11
0.33	-0.11	0.11	0.33	0.11	-0.11	0.33

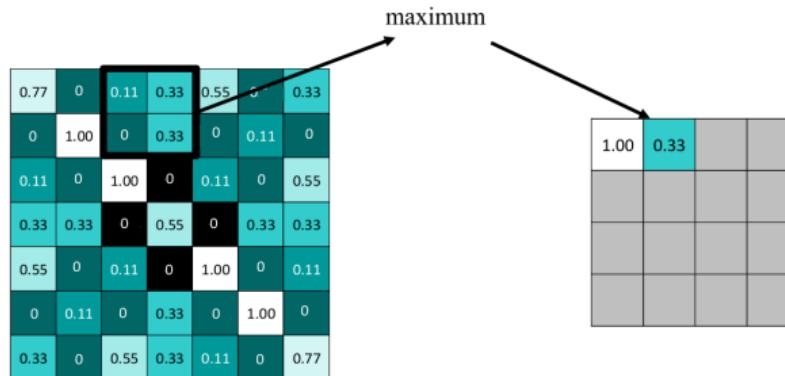


0.17	0	0.11	0.31	0.35	0	0.29
0	1.09	0	0.33	0	0.11	0
0.11	0	1.09	0	0.11	0	0.35
0.11	0.11	0	0.33	0	0.33	0.39
0.08	0	0.11	0	1.09	0	0.11
0	0.11	0	0.33	0	1.09	0
0.11	0	0.55	0.31	0.11	0	0.27
0.31	0	0.11	0	0.33	0	0.39
0.33	0	0.55	0	0.11	0	0.35
0.11	0	0.55	0	0.33	0	0.19
0.11	0	0.33	0	1.09	0	0.31
0.11	0	0.33	0	0.35	0	0.15
0	0.55	0	0.33	0	0.55	0
0.33	0	0.11	0	0.11	0	0.38
0.11	0	0.33	0.31	0.11	0	0.77
0	0.11	0	0.33	0	1.09	0
0.11	0	0.11	0	1.09	0	0.11
0.11	0.11	0	0.55	0	0.33	0.39
0.11	0	1.09	0	0.33	0	0.15
0.11	0	0.33	0	0.31	0	0.31
0.11	0	0.11	0	0.33	0	0.39

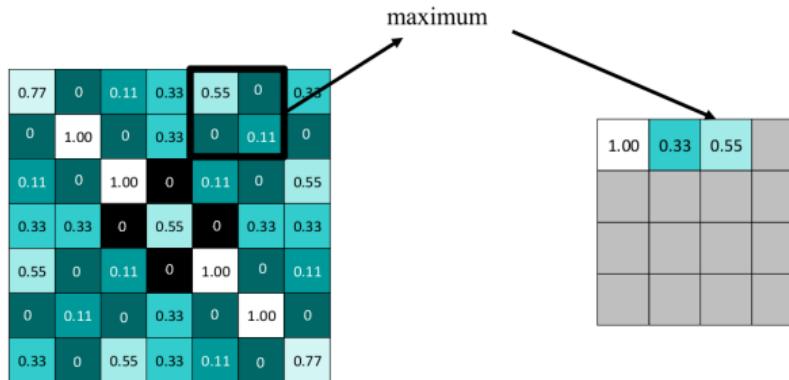
# Pooling: Shrinking the Image Stack



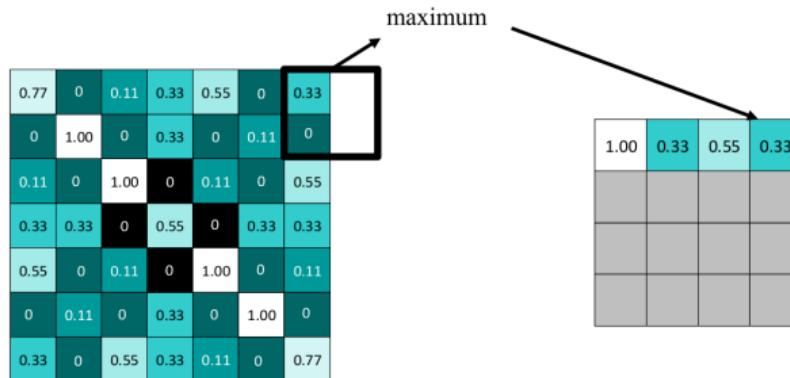
# Pooling: Shrinking the Image Stack



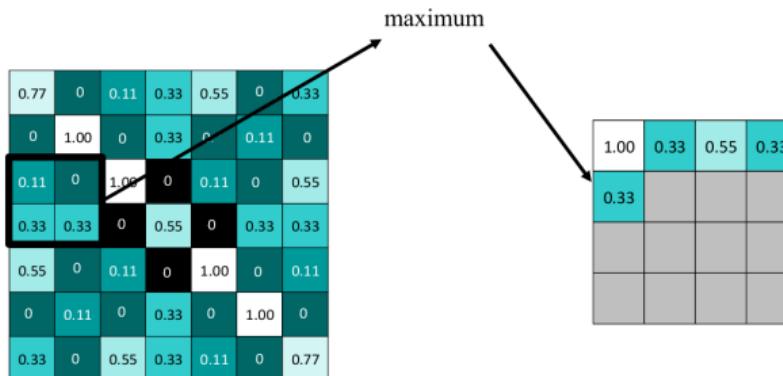
# Pooling: Shrinking the Image Stack



# Pooling: Shrinking the Image Stack



# Pooling: Shrinking the Image Stack



# Pooling: Shrinking the Image Stack

0.77	0	0.11	0.33	0.55	0	0.33
0	1.00	0	0.33	0	0.11	0
0.11	0	1.00	0	0.11	0	0.55
0.33	0.33	0	0.55	0	0.33	0.33
0.55	0	0.11	0	1.00	0	0.11
0	0.11	0	0.33	0	1.00	0
0.33	0	0.55	0.33	0.11	0	0.77

max pooling

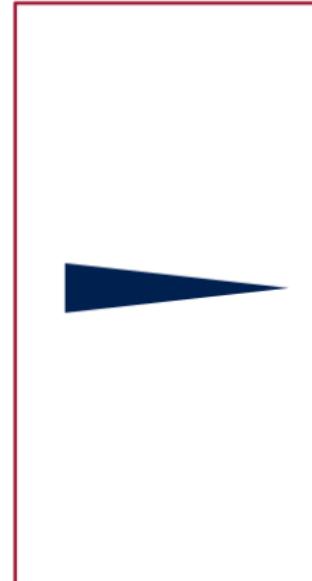
1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

# Repeat For All the Filtered Images

0.77	0	0.33	0.33	0.33	0	0.33
0	1.00	0	0.33	0	0.33	0
0.33	0	1.00	0	0.33	0	0.33
0.33	0.33	0	0.33	0	0.33	0.33
0.33	0	0.33	0	1.00	0	0.33
0.33	0.33	0	0.33	0	1.00	0
0.33	0	0.33	0.33	0	0	0.77

0.33	0	0.33	0	0.33	0	0.33
0	0.55	0	0.33	0	0.55	0
0.33	0	0.33	0	0.33	0	0.33
0	0.33	0	1.00	0	0.33	0
0.33	0	0.33	0	0.33	0	0.33
0	0.55	0	0.33	0	0.55	0
0.33	0	0.33	0	0.33	0	0.33

0.33	0	0.33	0.33	0.33	0	0.77
0	0.33	0	0.33	0	1.00	0
0.33	0	0.33	0	1.00	0	0.33
0.33	0.33	0	0.33	0	0.33	0.33
0.33	0	1.00	0	0.33	0	0.33
0	1.00	0	0.33	0	0.33	0
0.33	0	0.33	0.33	0.33	0	0.33



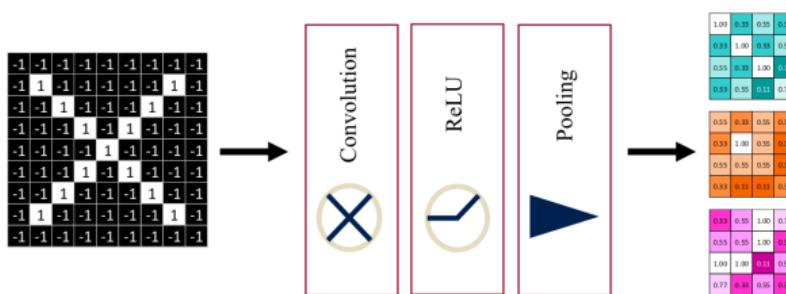
1.00	0.33	0.55	0.33
0.33	1.00	0.33	0.55
0.55	0.33	1.00	0.11
0.33	0.55	0.11	0.77

0.55	0.33	0.55	0.33
0.33	1.00	0.55	0.11
0.55	0.55	0.55	0.11
0.33	0.11	0.11	0.33

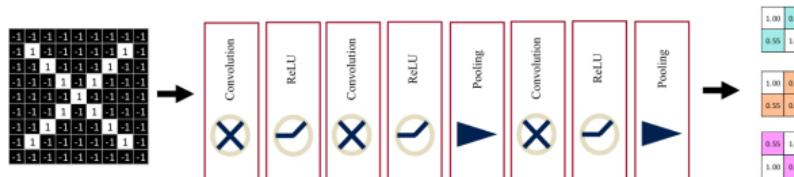
0.33	0.55	1.00	0.77
0.55	0.55	1.00	0.33
1.00	1.00	0.11	0.55
0.77	0.33	0.55	0.33

# Layers Get Stacked

- ▶ The output of one becomes the input of the next.

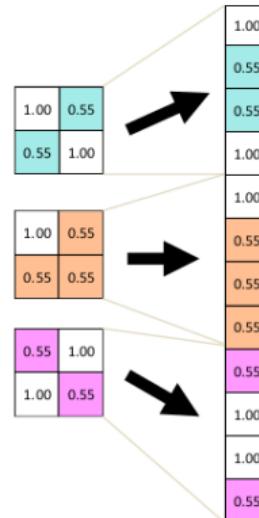


# Deep Stacking



# Fully Connected Layer

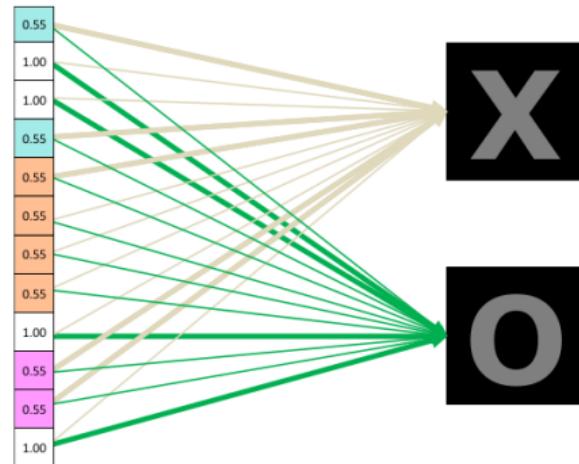
- ▶ Flattening the outputs before giving them to the **fully connected layer**.



# Fully Connected Layer



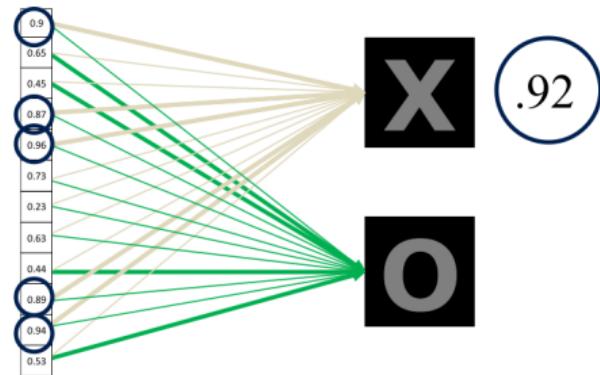
# Fully Connected Layer



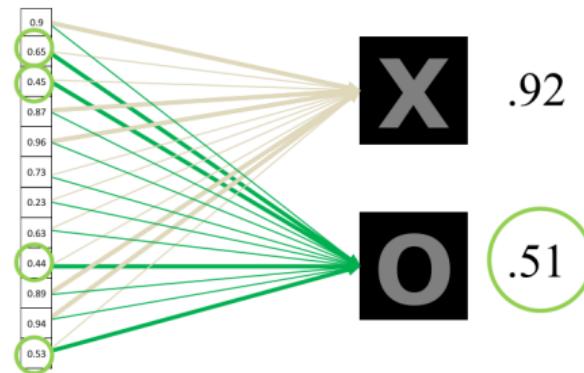
# Fully Connected Layer



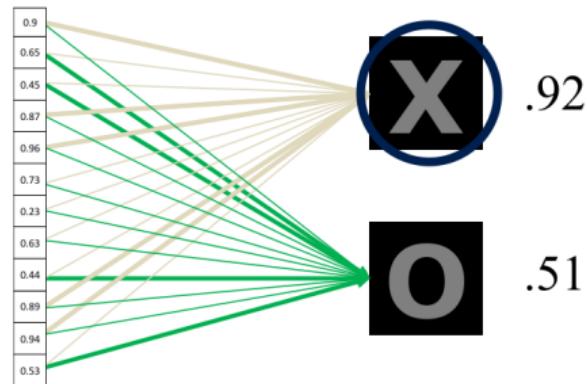
# Fully Connected Layer



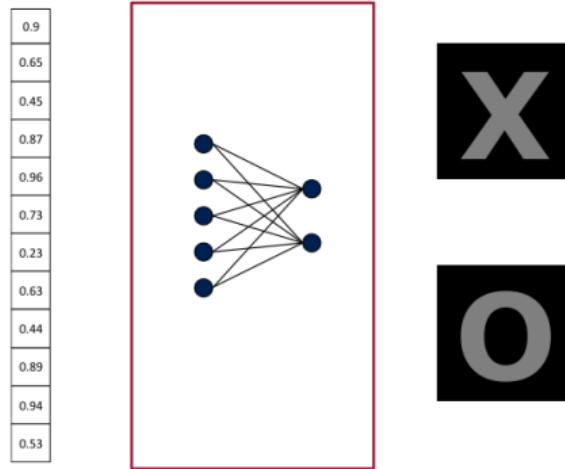
# Fully Connected Layer



# Fully Connected Layer

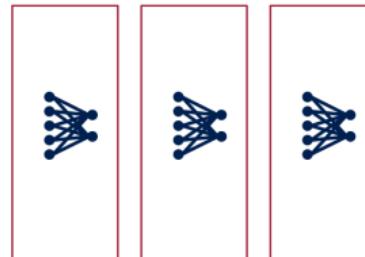


# Fully Connected Layer

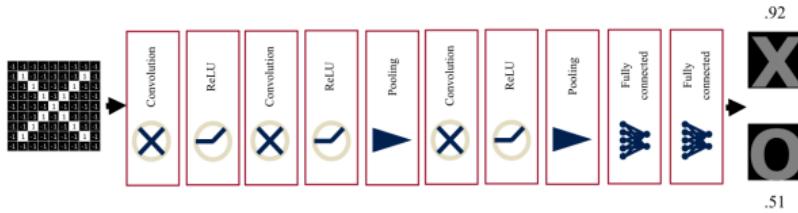


# Fully Connected Layer

0.9
0.65
0.45
0.87
0.96
0.73
0.23
0.61
0.44
0.89
0.94
0.53



# Putting It All Together







# CNN in TensorFlow



## CNN in TensorFlow (1/8)

- ▶ A **CNN** for the **MNIST** dataset with the following network.



## CNN in TensorFlow (1/8)

- ▶ A **CNN** for the **MNIST** dataset with the following network.
- ▶ Conv. layer 1: computes **32 feature maps** using a **5x5 filter** with ReLU activation.



## CNN in TensorFlow (1/8)

- ▶ A CNN for the MNIST dataset with the following network.
- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.
- ▶ Pooling layer 1: max pooling layer with a 2x2 filter and stride of 2.



## CNN in TensorFlow (1/8)

- ▶ A CNN for the MNIST dataset with the following network.
- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.
- ▶ Pooling layer 1: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Conv. layer 2: computes 64 feature maps using a 5x5 filter.



## CNN in TensorFlow (1/8)

- ▶ A CNN for the MNIST dataset with the following network.
- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.
- ▶ Pooling layer 1: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Conv. layer 2: computes 64 feature maps using a 5x5 filter.
- ▶ Pooling layer 2: max pooling layer with a 2x2 filter and stride of 2.



## CNN in TensorFlow (1/8)

- ▶ A CNN for the MNIST dataset with the following network.
- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.
- ▶ Pooling layer 1: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Conv. layer 2: computes 64 feature maps using a 5x5 filter.
- ▶ Pooling layer 2: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Dense layer: densely connected layer with 1024 neurons.



## CNN in TensorFlow (1/8)

- ▶ A CNN for the MNIST dataset with the following network.
- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.
- ▶ Pooling layer 1: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Conv. layer 2: computes 64 feature maps using a 5x5 filter.
- ▶ Pooling layer 2: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Dense layer: densely connected layer with 1024 neurons.
- ▶ Logits layer



## CNN in TensorFlow (2/8)

- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.



## CNN in TensorFlow (2/8)

- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.
- ▶ Input tensor shape: [batch\_size, 28, 28, 1]
- ▶ Output tensor shape: [batch\_size, 28, 28, 32]



## CNN in TensorFlow (2/8)

- ▶ Conv. layer 1: computes 32 feature maps using a 5x5 filter with ReLU activation.
- ▶ Input tensor shape: [batch\_size, 28, 28, 1]
- ▶ Output tensor shape: [batch\_size, 28, 28, 32]
- ▶ Padding same is added to preserve width and height.

```
# MNIST images are 28x28 pixels, and have one color channel
X = tf.placeholder(tf.float32, [None, 28, 28, 1])
y_true = tf.placeholder(tf.float32, [None, 10])

conv1 = tf.layers.conv2d(inputs=X, filters=32, kernel_size=[5, 5], padding="same",
activation=tf.nn.relu)
```



## CNN in TensorFlow (3/8)

- ▶ Pooling layer 1: max pooling layer with a 2x2 filter and stride of 2.



## CNN in TensorFlow (3/8)

- ▶ Pooling layer 1: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Input tensor shape: [batch\_size, 28, 28, 32]
- ▶ Output tensor shape: [batch\_size, 14, 14, 32]

```
pool1 = tf.layers.max_pooling2d(inputs=conv1, pool_size=[2, 2], strides=2)
```



## CNN in TensorFlow (4/8)

- ▶ Conv. layer 2: computes 64 feature maps using a 5x5 filter.



## CNN in TensorFlow (4/8)

- ▶ Conv. layer 2: computes 64 feature maps using a 5x5 filter.
- ▶ Input tensor shape: [batch\_size, 14, 14, 32]
- ▶ Output tensor shape: [batch\_size, 14, 14, 64]



## CNN in TensorFlow (4/8)

- ▶ Conv. layer 2: computes 64 feature maps using a 5x5 filter.
- ▶ Input tensor shape: [batch\_size, 14, 14, 32]
- ▶ Output tensor shape: [batch\_size, 14, 14, 64]
- ▶ Padding same is added to preserve width and height.

```
conv2 = tf.layers.conv2d(inputs=pool1, filters=64, kernel_size=[5, 5], padding="same",
activation=tf.nn.relu)
```



## CNN in TensorFlow (5/8)

- ▶ Pooling layer 2: max pooling layer with a 2x2 filter and stride of 2.



## CNN in TensorFlow (5/8)

- ▶ Pooling layer 2: max pooling layer with a 2x2 filter and stride of 2.
- ▶ Input tensor shape: [batch\_size, 14, 14, 64]
- ▶ Output tensor shape: [batch\_size, 7, 7, 64]

```
pool2 = tf.layers.max_pooling2d(inputs=conv2, pool_size=[2, 2], strides=2)
```



## CNN in TensorFlow (6/8)

- ▶ Flatten tensor into a batch of vectors.



## CNN in TensorFlow (6/8)

- ▶ **Flatten** tensor into a batch of vectors.
  - Input tensor shape: `[batch_size, 7, 7, 64]`
  - Output tensor shape: `[batch_size, 7 * 7 * 64]`

```
pool2_flat = tf.reshape(pool2, [-1, 7 * 7 * 64])
```



## CNN in TensorFlow (6/8)

- ▶ **Flatten** tensor into a batch of vectors.
  - Input tensor shape: `[batch_size, 7, 7, 64]`
  - Output tensor shape: `[batch_size, 7 * 7 * 64]`

```
pool2_flat = tf.reshape(pool2, [-1, 7 * 7 * 64])
```

- ▶ **Dense layer:** densely connected layer with **1024 neurons**.



## CNN in TensorFlow (6/8)

- ▶ **Flatten** tensor into a batch of vectors.
  - Input tensor shape: `[batch_size, 7, 7, 64]`
  - Output tensor shape: `[batch_size, 7 * 7 * 64]`

```
pool2_flat = tf.reshape(pool2, [-1, 7 * 7 * 64])
```

- ▶ **Dense layer:** densely connected layer with **1024 neurons**.
  - Input tensor shape: `[batch_size, 7 * 7 * 64]`
  - Output tensor shape: `[batch_size, 1024]`

```
dense = tf.layers.dense(inputs=pool2_flat, units=1024, activation=tf.nn.relu)
```



## CNN in TensorFlow (7/8)

- ▶ Add **dropout** operation; 0.6 probability that element will be kept

```
dropout = tf.layers.dropout(inputs=dense, rate=0.4)
```



## CNN in TensorFlow (7/8)

- ▶ Add **dropout** operation; 0.6 probability that element will be kept

```
dropout = tf.layers.dropout(inputs=dense, rate=0.4)
```

- ▶ **Logits layer**

- Input tensor shape: `[batch_size, 1024]`
- Output tensor shape: `[batch_size, 10]`

```
logits = tf.layers.dense(inputs=dropout, units=10)
```



## CNN in TensorFlow (8/8)

```
# define the cost and accuracy functions
cross_entropy = tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y_true)
cross_entropy = tf.reduce_mean(cross_entropy) * 100

# define the optimizer
lr = 0.003
optimizer = tf.train.AdamOptimizer(lr)
train_step = optimizer.minimize(cross_entropy)

# execute the model
init = tf.global_variables_initializer()

n_epochs = 2000
with tf.Session() as sess:
    sess.run(init)

    for i in range(n_epochs):
        batch_X, batch_y = mnist.train.next_batch(100)
        sess.run(train_step, feed_dict={X: batch_X, y_true: batch_y})
```





# Training CNNs

# Training CNN (1/4)

- ▶ Let's see how to use **backpropagation** on a **single convolutional layer**.

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



$h_{11}$	
$h_{21}$	$h_{22}$

# Training CNN (1/4)

- ▶ Let's see how to use **backpropagation** on a **single convolutional layer**.
- ▶ Assume we have an input  $X$  of size  **$3 \times 3$**  and a **single filter  $W$**  of size  **$2 \times 2$** .

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



$h_{11}$	
$h_{21}$	$h_{22}$

## Training CNN (1/4)

- ▶ Let's see how to use **backpropagation** on a **single convolutional layer**.
- ▶ Assume we have an input  $X$  of size  **$3 \times 3$**  and a **single filter  $W$**  of size  **$2 \times 2$** .
- ▶ **No padding** and **stride = 1**.

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



$h_{11}$	
$h_{21}$	$h_{22}$

# Training CNN (1/4)

- ▶ Let's see how to use **backpropagation** on a **single convolutional layer**.
- ▶ Assume we have an input  $X$  of size  **$3 \times 3$**  and a **single filter  $W$**  of size  **$2 \times 2$** .
- ▶ **No padding** and **stride = 1**.
- ▶ It generates an **output  $H$**  of size  **$2 \times 2$** .

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



$h_{11}$	
$h_{21}$	$h_{22}$



## Training CNN (2/4)

- ▶ Forward pass

## Training CNN (2/4)

### ► Forward pass

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



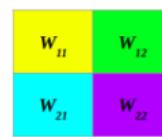
$h_{11}$	
$h_{21}$	$h_{22}$

$$h_{11} = W_{11}X_{11} + W_{12}X_{12} + W_{21}X_{21} + W_{22}X_{22}$$

## Training CNN (2/4)

### ► Forward pass

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



$h_{11}$	
$h_{12}$	
$h_{21}$	$h_{22}$

$$h_{11} = W_{11}X_{11} + W_{12}X_{12} + W_{21}X_{21} + W_{22}X_{22}$$

$$h_{12} = W_{11}X_{12} + W_{12}X_{13} + W_{21}X_{22} + W_{22}X_{23}$$

## Training CNN (2/4)

### ► Forward pass

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



$h_{11}$	
	$h_{12}$
$h_{21}$	
	$h_{22}$

$$h_{11} = W_{11}X_{11} + W_{12}X_{12} + W_{21}X_{21} + W_{22}X_{22}$$

$$h_{12} = W_{11}X_{12} + W_{12}X_{13} + W_{21}X_{22} + W_{22}X_{23}$$

$$h_{21} = W_{11}X_{21} + W_{12}X_{22} + W_{21}X_{31} + W_{22}X_{32}$$

## Training CNN (2/4)

### ► Forward pass

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$



$h_{11}$	
	$h_{12}$
$h_{21}$	
	$h_{22}$

$$h_{11} = W_{11}X_{11} + W_{12}X_{12} + W_{21}X_{21} + W_{22}X_{22}$$

$$h_{12} = W_{11}X_{12} + W_{12}X_{13} + W_{21}X_{22} + W_{22}X_{23}$$

$$h_{21} = W_{11}X_{21} + W_{12}X_{22} + W_{21}X_{31} + W_{22}X_{32}$$

$$h_{22} = W_{11}X_{22} + W_{12}X_{23} + W_{21}X_{32} + W_{22}X_{33}$$

## Training CNN (3/4)

- ▶ Backward pass
- ▶ E is the error:  $E = E_{h_{11}} + E_{h_{12}} + E_{h_{21}} + E_{h_{22}}$

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$

$W_{11}$	$W_{12}$
$W_{21}$	$W_{22}$

$h_{11}$	$h_{12}$
$h_{21}$	$h_{22}$

# Training CNN (3/4)

- ▶ Backward pass
- ▶ E is the error:  $E = E_{h_{11}} + E_{h_{12}} + E_{h_{21}} + E_{h_{22}}$

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$

$W_{11}$	$W_{12}$
$W_{21}$	$W_{22}$

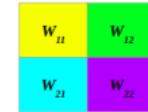
$h_{11}$	$h_{12}$
$h_{21}$	$h_{22}$

$$\frac{\partial E}{\partial W_{11}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{11}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{11}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{11}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{11}}$$

# Training CNN (3/4)

- ▶ Backward pass
- ▶ E is the error:  $E = E_{h_{11}} + E_{h_{12}} + E_{h_{21}} + E_{h_{22}}$

$X_u$	$X_{i_2}$	$X_{i_3}$
$X_{i_1}$	$X_{i_2}$	$X_{i_3}$
$X_u$	$X_{i_2}$	$X_{i_3}$



$h_{ii}$	$h_{i2}$
$h_{21}$	$h_{22}$

$$\frac{\partial E}{\partial W_{11}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{11}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{11}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{11}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{11}}$$

$$\frac{\partial E}{\partial W_{12}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{12}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{12}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{12}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{12}}$$

# Training CNN (3/4)

- ▶ Backward pass
- ▶ E is the error:  $E = E_{h_{11}} + E_{h_{12}} + E_{h_{21}} + E_{h_{22}}$

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$

$w_{11}$	$w_{12}$
$w_{21}$	$w_{22}$

$h_{11}$	$h_{12}$
$h_{21}$	$h_{22}$

$$\frac{\partial E}{\partial W_{11}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{11}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{11}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{11}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{11}}$$

$$\frac{\partial E}{\partial W_{12}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{12}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{12}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{12}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{12}}$$

$$\frac{\partial E}{\partial W_{21}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{21}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{21}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{21}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{21}}$$

# Training CNN (3/4)

- ▶ Backward pass
- ▶ E is the error:  $E = E_{h_{11}} + E_{h_{12}} + E_{h_{21}} + E_{h_{22}}$

$X_{11}$	$X_{12}$	$X_{13}$
$X_{21}$	$X_{22}$	$X_{23}$
$X_{31}$	$X_{32}$	$X_{33}$

$w_{11}$	$w_{12}$
$w_{21}$	$w_{22}$

$h_{11}$	$h_{12}$
$h_{21}$	$h_{22}$

$$\frac{\partial E}{\partial W_{11}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{11}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{11}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{11}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{11}}$$

$$\frac{\partial E}{\partial W_{12}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{12}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{12}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{12}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{12}}$$

$$\frac{\partial E}{\partial W_{21}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{21}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{21}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{21}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{21}}$$

$$\frac{\partial E}{\partial W_{22}} = \frac{\partial E_{h_{11}}}{\partial h_{11}} \frac{\partial h_{11}}{\partial W_{22}} + \frac{\partial E_{h_{12}}}{\partial h_{12}} \frac{\partial h_{12}}{\partial W_{22}} + \frac{\partial E_{h_{21}}}{\partial h_{21}} \frac{\partial h_{21}}{\partial W_{22}} + \frac{\partial E_{h_{22}}}{\partial h_{22}} \frac{\partial h_{22}}{\partial W_{22}}$$

# Training CNN (4/4)

- ▶ Update the weights  $W$



$$W_{11}^{(\text{next})} = W_{11} - \eta \frac{\partial E}{\partial W_{11}}$$

$$W_{12}^{(\text{next})} = W_{12} - \eta \frac{\partial E}{\partial W_{12}}$$

$$W_{21}^{(\text{next})} = W_{21} - \eta \frac{\partial E}{\partial W_{21}}$$

$$W_{22}^{(\text{next})} = W_{22} - \eta \frac{\partial E}{\partial W_{22}}$$



# Summary



# Summary

- ▶ Receptive fields and filters
- ▶ Convolution operation
- ▶ Padding and strides
- ▶ Pooling layer
- ▶ Flattening, dropout, dense



## Reference

- ▶ Tensorflow and Deep Learning without a PhD  
<https://codelabs.developers.google.com/codelabs/cloud-tensorflow-mnist>
- ▶ Ian Goodfellow et al., Deep Learning (Ch. 9)
- ▶ Aurélien Géron, Hands-On Machine Learning (Ch. 13)



# Questions?